

Likelihood approximations distort the ion channel kinetic information in macroscopic currents

Luciano Moffatt^{1*}

*For correspondence:
lmoffatt@qi.fcen.uba.ar

¹Instituto de Química Física de los Materiales, Medio Ambiente y Energía (INQUIMAE), CONICET; Facultad de Ciencias Exactas y Naturales, Universidad de Buenos Aires, Ciudad de Buenos Aires, C1428EHA, Argentina

Abstract Kinetic schemes are usually fitted to macroscopic ion-channel currents by least squares on the deterministic mean, leaving the gating fluctuations unmodelled. Likelihoods that use them, including Bayesian recursive filters, are two decades old and little used: a filter tracks the data even when the uncertainty it reports is wrong. In our own comparison of nine gating schemes for the P2X₂ receptor, dropping the recursion changed the ranking. The test is classical: at the true parameters the score, the log-likelihood's derivative, averages to zero, and its covariance across recordings equals the Fisher information it reports. We measure bias and distortion along a ladder from least squares to a filter conditioned on both ends of each acquisition interval, for a two-state scheme over channel number, noise and acquisition interval, against exact simulations of the process they approximate. The boundary-conditioned filter is calibrated over most of the sweep, departing toward few channels and low noise, where the Gaussian moment closure on the occupancy is misspecified. Least squares becomes calibrated once instrumental noise exceeds gating noise, which happens at an instrumental variance proportional to channel number. No region of the measured plane gives both information about the unitary current and a trustworthy least-squares error bar, the two boundaries one to three decades of noise apart.

Introduction

A channel opens and shuts at random, and how long it stays in each conformation is the elementary random quantity of gating. What an experiment sees of it depends on how many channels contribute. With one channel in the patch the durations are read off the record and the rate constants come from their distributions (*Neher and Stevens, 1977; Colquhoun and Hawkes, 1981*). In the currents that carry an action potential so many channels contribute that the durations average away, the current follows the deterministic relaxation of the Hodgkin and Huxley description (*Hodgkin and Huxley, 1952*), and rate constants are obtained by fitting that relaxation, in current practice by least squares.

Most macroscopic patch-clamp recordings are in neither place. A few hundred to a few thousand channels contribute, individual openings cannot be resolved, and the averaging is incomplete. The durations are still there, and what they leave behind is a fluctuation around the mean current that carries the memory of the gating process. Samples taken closer together than the channel's relaxation time are correlated, so a sample taken faster than that time partly repeats what the one before it already said.

That fluctuation carries what the mean cannot. Channel number and unitary current enter the mean current as a product, so no fit to the mean separates them. The standard analysis discards the fluctuation anyway: the four approaches surveyed by *Clerx et al. (2019)* are deterministic fits with a root-mean-square objective, and the IonBench benchmark states that it does not cover optimisation approaches for stochastic models (*Owen and Mirams, 2025*). Discarding it usually leaves the point estimate standing. What it damages is the count of independent residuals, the number that turns a residual sum of squares into a confidence interval, so a fit that treats correlated residuals as independent behaves as though the recording had supplied more observations than it did. *Del Core and Mirams (2025)* state the same cost in the field's own terms, that standard methods built for deterministic models do not distinguish stochastic channel gating from measurement noise and return biased estimates.

A Markov description models the dependence rather than assuming it away. A channel with a small number of conformational states, connected by rate constants that hold over the whole record, generates the mean current and the covariance of the fluctuations from the same constants, so the temporal structure of the noise becomes a second measurement of the same parameters. Those constants also mean something outside the record. They are properties of the protein, comparable across conditions, preparations and laboratories, and they connect what is measured to a mechanism. A correlation function fitted freely to the residuals delivers neither of those.

Kinetic information has been read from that correlation since the 1970s, first through stationary power spectra (*Katz and Miledi, 1970, 1972; Anderson and Stevens, 1973*), from which *Anderson and Stevens (1973)* obtained a unitary conductance and a closing rate jointly with standard errors on both, and then through the exact non-stationary two-time covariance of a population of identical independent channels, estimated from ensembles of 256 to 504 repeated sweeps (*Conti et al., 1980; Sigworth, 1981*). The whole-record likelihoods that followed differ along two axes and are best read that way. The first is how much of the gating stochasticity the likelihood models: nothing, then the variance at each sample (*Milescu et al., 2005*), then its correlation in time, whether by carrying the full two-time covariance at a cost that keeps it to a hundred or two selected points (*Celentano and Hawkes, 2004*), or by propagating the occupancy recursively and conditioning it on each new datum at a cost linear in the record (*Moffatt, 2007; Stepanyuk et al., 2011; Münch et al., 2022*). The second axis is the observable itself. An acquisition system reports an average over each interval rather than an instantaneous sample, and the most recent member conditions that interval average on the channel states at both of its ends (*Moffatt and Pierdominici-Sottile, 2025*).

None of these has displaced least squares. Cost is part of it, and *Del Core and Mirams (2025)* make that case. There is also something these methods ask an analyst to give up. A least-squares fit predicts the current at every point from the parameters alone, so the residual can be read directly, and read twice. Its amplitude reports the quality of the fit against what the gating variance and the baseline noise should deliver, and its correlation in time reports what the fit leaves out. Both readings appeal to knowledge from outside the fit, which is what makes them checks.

Modelling the correlation is what takes those readings away. A recursive likelihood predicts each interval from the parameters and from the data already observed, so the previous sample enters the prediction of the next one and the predicted trace follows the record by construction. It follows the record when the scheme is right, when the scheme is wrong, and when the code has a mistake in it. The limiting case of the argument carries no kinetics at all and states that each sample equals the last one plus noise: where the interval is short against the relaxation time and the instrumental noise is small, it predicts almost perfectly and says nothing about the channel. Such a likelihood also standardises its residual against its own predicted variance, so the amplitude reading becomes a statement of internal consistency, and the correlation reading is spent because the recursion subtracts at each step what it has just seen. The same blindness covers two questions that an experimenter cannot separate and has no reason to: whether the approximation

is adequate, and whether the code that computes it is correct.

That is not a hypothetical. Nine kinetic schemes for the P2X₂ receptor were ranked by Bayesian evidence twice over the same outside-out patch recordings, in work of our own, and the two rankings were not the same (*Moffatt and Pierdominici-Sottile, 2025*). What separated them was the recursion. Both likelihoods averaged the current over the acquisition interval and both carried the same gating variance; one conditioned each interval on the data already observed and the other did not. That the ranking moved with the approximation was reported. Which of the two orderings to believe was not, because it turns on how far each likelihood's reported information sits from the information it delivers, and that had never been measured for either.

The apparatus that answers this is classical. For a correctly specified likelihood the score, the gradient of the log-likelihood with respect to the parameters, has mean zero at the parameters that generated the data, and its covariance across independent recordings equals the Fisher information the likelihood reports (*Huber, 1967; White, 1982*). An approximation satisfies neither statement at those parameters, and the two departures are measurable quantities rather than tests to be passed. Goodness of fit does not answer the question, since an approximation can reproduce the mean current closely and still report an information matrix wrong by more than an order of magnitude, and comparing two approximations against each other does not answer it either, because both may be wrong. Measuring the departures takes an ensemble of recordings drawn from known parameters, which is why this is a simulation study and not something a reader can run on a single record. What makes it possible at all for this class of model is an asymmetry between simulation and inference: the forward process can be simulated exactly while its likelihood cannot be computed exactly, so the simulator supplies ground truth for what the likelihood approximates.

Two alternatives to a mechanistic likelihood are worth placing here, since both are in use. The standard statistical repair for correlated residuals is generalized least squares with autoregressive or moving-average errors, which whitens the residual sequence with a correlation function estimated from the residuals themselves, and *Lei et al. (2020)* have run exactly that on ion-channel data. Gating puts that family out of reach for structural reasons. An autoregressive moving-average error model is stationary by construction, one covariance function for the whole trace, while the gating covariance tracks the mean current, vanishes when the current does, and restarts at every concentration or voltage jump. It contains no channel number and no unitary current, so it cannot return the two parameters the fluctuations are richest in. Its timescales are free parameters, where the Markov model ties them to the eigenvalues of the rate matrix that already fixed the mean, so free timescales absorb the information the mechanism would have contributed. *Lei et al. (2020)* report that the autoregressive model widens the posterior without predicting better. The other alternative is non-stationary fluctuation analysis, the one fluctuation-based method in routine use, which regresses the variance of the current against its mean and returns the unitary current, the channel number and the peak open probability (*Stepanyuk et al., 2014*). This paper does not compete with it over how much amplitude information a macroscopic recording holds, never benchmarks against it, and no comparison with the variance-mean regression should be read into any result below.

Because every member is an approximation, the measurement returns a domain rather than a verdict. The question is where a likelihood reports honestly, and the answer is a map whose boundaries are level sets of a continuous diagnostic rather than thresholds. Position on it is set by the number of channels, by the acquisition interval in units of the channel's relaxation time, and by the instrumental noise in units of the gating noise, held at fixed spectral density so that shortening the interval raises the noise in each sample and leaves the noise per unit bandwidth unchanged. The measurements are made on a two-state scheme at an open probability of one half, driven by a single concentration jump, compared at the level of the likelihood rather than of a fitted experiment. The small model isolates the error the approximation itself commits, with no confounding from a misspecified kinetic scheme, it keeps the exact simulation cheap enough to

run thousands of replicates at every design cell, and it is the most favourable case for a Gaussian treatment of the occupancy and therefore a lower bound on the error for anything richer. Open probability, stationary protocols and larger schemes are left out for reasons that follow from the same design, and are stated with the results.

Two things come out of the map, and they answer different readers. The first is where each approximation misreports what the data determine, including the most complete one, since a method whose limits are unmapped has been endorsed rather than characterised. The second is the region in which a classical fit reports an interval that matches the spread it delivers, and it carries an ordering that is falsifiable and does not depend on where the level sets are placed: there is no region of the measured plane in which the fluctuations still carry information about the unitary current or the channel number and the classical interval is simultaneously trustworthy, the two boundaries being separated by one to three decades of instrumental noise everywhere they were measured.

Against that history, what is new here is narrow and is worth stating narrowly. A joint estimate of rates and amplitudes from fluctuation data, carrying standard errors, is **Anderson and Stevens (1973)**. What this paper adds, to a non-stationary protocol read from a single record, is the third element of the combination: an interval whose calibration has been verified rather than assumed. The thousands of replicate recordings used below certify the method; they are not what applying it requires. Separately from the map, there is a reading an experimenter can take from a recording they already have. The integrated autocorrelation of the standardized least-squares residual counts how many intervals the fit effectively has for every interval it believes it has, and it costs nothing, since the fit was going to be made anyway. Two design variables that look as though they should matter, the channel count and the instrumental noise, enter the answer through that one statistic and through nothing else. The acquisition interval does not, and the reader knows it, because they chose it when they set the sampling rate, as they did the analogue filter, which contributes to the same statistic on a record sampled faster than its own bandwidth. Those two numbers together say what a classical fit is doing to the width of the interval it reports, in units a reader can act on: a distortion of 1.15 is a 7% error on a confidence interval, 2 is 41%, and 4 turns a nominal 95% interval into about 68%.

The macroscopic interval likelihood and its approximations

A macroscopic current is the summed signal of a population of ion channels, each of which behaves as an independent continuous-time Markov chain over a small set of conformational states. We write N_{ch} for the number of channels and K for the number of kinetic states of one channel ($K = 2$ throughout this paper, carried symbolically where it costs nothing). Each channel is governed by a generator (rate) matrix $\mathbf{Q} \in \mathbb{R}^{K \times K}$, whose transition matrix over an elapsed time t is the row-stochastic $\mathbf{P}(t) = e^{\mathbf{Q}t}$ with entries $P_{i \rightarrow j}(t)$. A channel in state k carries a conductance γ_k , and we collect these in the vector $\boldsymbol{\gamma} \in \mathbb{R}^K$.

Three marks, and each one says one thing

An arrow means the object spans the acquisition window. $P_{i \rightarrow j}(\Delta)$ is a transition across it, $\bar{y}_{i \rightarrow j}$ an average over it, $\bar{y}_{0 \rightarrow \Delta}$ a recorded value formed from it. **No arrow means the object lives at one instant:** μ_0 and Σ_0 are the occupancy mean and covariance where the window opens, μ_Δ and Σ_Δ where it closes. **A dot in a slot means that end has been collapsed**, as in $\bar{y}_{i \rightarrow \cdot}$. For a quantity conditioned on the pair the dot denotes the P -weighted collapse, $a_{i \rightarrow \cdot} := \sum_j P_{i \rightarrow j}(\Delta) a_{i \rightarrow j} = \mathbb{E}[a \mid X_0 = i]$, so a collapsed object is defined the moment its pair-resolved version is. For a probability the weight is already inside and the collapse is the plain margin.

Letters name states and the arrow names the time relation. Collapsing one end does not remove the arrow, because the object still spans: $\bar{y}_{i \rightarrow \cdot}$ is an average over the window conditioned on one of its ends, not a quantity at an instant.

The arrow says what the object is, not how it was obtained. Σ_Δ carries no arrow even though a propagator was used to reach it, because it is the covariance of the occupancy at one time with itself.

A quantity indexed by pairs is written once and read according to the operator it meets, as a $K \times K$ matrix under an entrywise product and as a vector of length K^2 under a contraction against $\Sigma_{0 \rightarrow \Delta}$, the two being the same array. Throughout, $\circ$ is the entrywise product and $\mathbf{1}$ the vector of ones. Where a whole recording is in view rather than one window, a subscript i indexes the window, as in $\bar{y}_i$; inside a window the two ends are 0 and Δ .

The likelihood we need is the probability of a recorded sample given the kinetic parameters and the data seen so far. Writing it exactly calls for two separate objects. The first is a microscopic description of the population, one that keeps each channel as a discrete entity and carries the state of the ensemble as an integer occupancy vector. That description is explicit, and the recursion over it is the microscopic filter of *Moffatt (2007)*, but the number of occupancy vectors is $\binom{N_{\text{ch}}+K-1}{K-1}$, which is $N_{\text{ch}} + 1$ at $K = 2$, and every sample costs the square of that, which puts it out of reach at the channel counts of a macroscopic recording. The second is the law of the conductance averaged over the acquisition window, conditional on that description, and this one has no closed expression at any channel count, although its low-order moments do. The two obstructions are of different kinds: the first is cost, the second is that the object is not available in closed form at all. Every likelihood that runs at macroscopic channel counts replaces both objects with Gaussians, and the members of the family studied here differ in which object the second Gaussian describes, in what each is conditioned on and in how its variance is accounted. The rest of this section unpacks that sentence.

The observable is an interval average

A recorded sample is the current averaged over the acquisition window rather than sampled at an instant. The current is low-pass filtered in analogue form before it reaches the converter, so each stored value carries a weighted average of the recent current and not a reading of it at a point. Writing the sampling period as Δ and taking the weighting to be uniform across it, the quantity a likelihood must describe is the windowed average

$$\bar{y}_i = \frac{1}{\Delta} \int_{t-\Delta}^t y(\tau) d\tau + \eta_i, \quad (1)$$

where $y(\tau)$ is the instantaneous population current and η_i is the instrumental noise accumulated over the window. This interval average is the physical observable. When Δ is much shorter than the fastest relaxation of $\mathbf{Q}$ the average sits close to the instantaneous current and the distinction is negligible; when Δ is comparable to or longer than that relaxation the two diverge, and a likelihood built on instantaneous sampling describes the wrong random variable.

Uniform weighting is a first approximation to what the electronics do. The analogue stage of a patch-clamp amplifier is normally a Bessel low-pass, chosen because its group delay is nearly flat and its step response barely overshoots, so the effective measurement kernel is that filter's impulse response, a smooth shape whose tails reach into the neighbouring intervals. The uniform window keeps the entire kernel inside one acquisition interval, which is what lets a likelihood conditioned on the two ends of that interval account for the averaging without reference to anything outside it. A Bessel kernel breaks that containment twice over: it correlates successive recorded values through the signal, and it colours the instrumental noise that Eq. 1 takes to be white. We keep the uniform window throughout, in the simulator as well as in every likelihood (Methods), so that data and likelihood describe the same observable and the closures below are the approximation the diagnostic sees. The uniform window is itself an approximation to the impulse response of the anti-aliasing filter of a real recording, and the Discussion prices it; it does not enter the measurement here because the simulator and every member share it. Carrying a general kernel, with the coloured noise that comes with it, through the same construction is left to later work.

Two Gaussian approximations

Two closures are available, on two different objects. Every member of the family makes the first, and every member ends up with a Gaussian for the recorded sample; what separates them is whether that Gaussian costs a second approximation. It does not for a member that predicts the current at an instant, because the instantaneous population current is γ contracted against the occupancy vector at that instant, a linear function of it, and the linear image of a Gaussian is Gaussian. It does for a member that predicts the interval average, which is a functional of the path taken inside the window and not a function of the occupancies at any one time. To close a distribution is to replace it, at that step, by the Gaussian with the same mean and covariance, discarding every higher moment. The word is borrowed from moment hierarchies, where that substitution lets a finite set of equations stand for an infinite one, and it fits here for the same reason: conditioning the occupancies on an observation takes them out of the multinomial family, and the conditional distribution after t samples belongs to no fixed-dimensional family that survives both propagation and conditioning, so carrying it exactly means carrying one probability per occupancy vector, a count that grows with N_{ch} , whereas (μ, Σ) is the same size at every step and at every channel count. Matching moments at each step is not the same as keeping them true across a recording: the update a closure licenses is itself approximate, so the pair carried forward drifts from the pair the exact filter would carry.

Occupancy closure. *Exact object:* one probability per occupancy vector, the microscopic description above, of a size that grows with N_{ch} . *Replacement:* a Gaussian carrying the per-channel pair (μ, Σ) , the same size at every channel count; because the occupancies sum to one, Σ is singular and the Gaussian lives on the simplex. *Validity:* two conditions, not one. The channel count must be high enough for the counts to behave as a continuum, which is the central limit theorem and fails as N_{ch} falls; and the recorded sample must not resolve the step in predicted current made by moving one channel between states, which fails as the instrumental noise falls. The second condition is why the boundary of the failing region in the Results runs across the plane of channel count against noise instead of down a line of constant N_{ch} . *What is closed* depends on whether the occupancies are conditioned on the data. A member that only propagates them holds them exactly multinomial at every sample time, since independent channels start multinomial and propagation preserves the family, so its two carried moments are the true ones and the closure discards shape alone. A member that updates them on each sample leaves the multinomial family at the first update and never returns to it, and its carried moments are approximate from then on.

Interval-signal closure. *Exact object:* the law of the conductance averaged over one acquisition window, a telegraph-like object (below) with no closed form; its low-order moments are available in closed form and are assembled from single-channel quantities (below). *Replacement:* a Gaussian carrying those moments. *Validity:* the window must span many transitions, or the instrumental noise must be large enough to make the departure immaterial. The second half of that is exact rather than heuristic: noise independent of the gating current leaves every cumulant of the recorded sample above the second unchanged while adding to its variance, so the standardized departures shrink as the noise grows. *What is closed* depends on how much of the window the conductance is conditioned on, which is the ladder below. A member that predicts the current at an instant does not make this second closure, and the statement is about which object gets closed. Such a member does give the recorded sample a Gaussian: the instantaneous population current is γ applied to the occupancies at a single time, so it is a linear image of the occupancy Gaussian and is Gaussian already, at no further cost in moments. What it drops is the evolution of the channels inside the window, one snapshot standing in for the average the recording holds, which misspecifies the random variable rather than closing it. The substitution also runs in the unhelpful direction, since averaging over the window is itself what brings the law closer to Gaussian.

The two are controlled by different quantities, the occupancy closure by the channel count and the interval-signal closure by the window measured against the gating timescale, while instrumental noise relaxes both by making the recorded sample more Gaussian. That is what makes them separable by design, and it is why the Results sweep a plane of channel count against noise at one window and then move the window along an axis of its own.

Telegraph-like names the following object. A single channel switching among its conductance levels traces a piecewise-constant path, the random telegraph signal when $K = 2$, and what the window delivers is its integral over one acquisition interval,

$$A_{0 \rightarrow \Delta} = \int_0^\Delta \gamma(X_s) ds, \quad (2)$$

where X_s is the state of the channel at time s . Its law is mixed: an atom at $\Delta\gamma_i$ holding the probability that the channel never left state i , and a continuous part from the paths that did switch. The moments, on the other hand, are cheap. Conditional on the states at the two ends of the window, the moment generating function of $A_{0 \rightarrow \Delta}$ is one entry of the tilted semigroup $e^{(Q + \alpha\Gamma)\Delta}$ divided by

the transition probability, with $\Gamma = \text{diag}(\gamma)$, so every moment is one derivative away and costs one matrix exponential, at any K .

The density is another matter, and it is where K enters: a series over the number of transitions inside the window, exact at every K but closed only at $K = 2$, with the terms proliferating from $K = 3$ on. Having few distinct conductances does not rescue it, which is worth saying because channel models are exactly of that kind: every shut state carries $\gamma = 0$ and the open states often share one γ , so $A_{0 \rightarrow \Delta}$ is the open conductance times the time spent in the open aggregate however large K is. That aggregate is Markov only if $\mathbf{Q}$ is lumpable, which for the schemes that motivate this work it is not (*Colquhoun and Hawkes, 1981*). The population multiplies the difficulty rather than adding to it: the recorded current sums over channels, so the laws convolve and the pieces proliferate in N_{ch} as well, and a recorded sample constrains only the sum, so the posterior over channels does not factorize (*Moffatt, 2007*) and an exact recursive filter has to carry a distribution over joint occupancy vectors. The closure is indifferent to all of it. Replacing $A_{0 \rightarrow \Delta}$ by a Gaussian truncates the transform at the second cumulant, and the moments that leaves it needing are single-channel objects that do not depend on N_{ch} or on the topology, delivered by two derivatives of that same $K \times K$ semigroup (below). That asymmetry, an exact law that degrades with K and with N_{ch} against moments that do neither, is what justifies the closure beyond the two-state case run here.

The higher moments these closures discard do not vanish. They reappear in the data as temporal correlation that the likelihood does not predict, which is what the correlation term of the information-distortion decomposition measures (Diagnostics). Three regimes follow, and we name them here so the maps in the Results have vocabulary: the *microscopic* regime (few channels, where the occupancy closure fails and only the microscopic description is exact), the *telegraphic* regime (windows far shorter than the relaxation time, where the interval-signal closure fails), and the *Gaussian* regime, where both closures hold because the channel count is high enough for the occupancy Gaussian while the window is long enough, or the instrumental noise large enough, for the interval-signal Gaussian. Any Gaussian likelihood of this kind must fail somewhere. The aim of this paper is to show that the failures are the predicted degradation of these two closures rather than defects of a particular algorithm.

Least squares, and the two axes of the family

Before any of the choices above comes a prior question: does the method model the gating fluctuations at all? Classical nonlinear least squares answers no. It fits the deterministic mean current and assigns one constant variance to every residual, so it carries neither an occupancy distribution nor a conductance conditioned on anything, and the temporal structure the two closures approximate is absent from it. Least squares is therefore not a member of the family. It is the bottom rung of the ladder of computational cost this paper walks, and it is the anchor against which every member above it is measured, because it is the method most readers are using.

The methods differ along two axes. One axis is the window: how much of the acquisition interval the predicted current is conditioned on. The other is recursion: whether the occupancy distribution is conditioned on the data as the recording proceeds.

The window axis is not the family's property alone. What the amplifier delivers is an average over the acquisition interval, so any method that predicts a recorded value has to decide whether to predict the current at an instant or the current averaged over the window, and that decision is available at every level of the gating description, least squares included. Least squares therefore has two arms here. They carry one constant variance and are otherwise identical, so the pair isolates the acquisition average with nothing else moving, which no pair inside the family does.

Along the recursion axis, a recursive (prequential) member factorizes the likelihood as $\prod_t q_t(\bar{y}_t \mid \bar{y}_{1:t-1})$, so the occupancy distribution carried into each window is the posterior left by the previous windows and the hidden state is conditioned on the data. A non-recursive (open-loop) member factorizes as $\prod_t m_t(\bar{y}_t)$, with each window's predictive propagated from the initial condition and never conditioned on the observed trace.

The members are named from the two axes and the names are read off them. The suffix is the recursion axis: NR for a non-recursive member, R for a recursive one. The prefix is the window axis, and it says how much of the interval the conductance is conditioned on: no prefix for the current at an instant, M for the interval mean given the state at the window's start, I for the interval mean given the states at both of its ends. Five members follow rather than six, because without an update the state at the end of a window is never observed and the two interval prefixes describe the same method on the non-recursive side: NR and R predict an instant, INR is the non-recursive interval member, and MR and IR are where the start state and the boundary separate. The M is the mean conductance and not an abbreviation of macroscopic; every member of the family is macroscopic. A third choice, open only to the members that stop at the start state, is the form given to the interval variance, and it accounts for the one letter that names neither axis: VR is MR with the residual interval variance in place of the total one, the two differing as set out at Eq. 26 below. Least squares carries no gating description at all and sits off this lattice, and we write LSE for the fit to the instantaneous deterministic mean and ILSE for the fit to the same mean averaged over the interval, where the I records the window setting and nothing else. Eight in all: five on the lattice, VR beside MR, and the two least-squares arms. The window decision is what separates LSE from ILSE, NR from INR, and R from MR and IR.

These are the two closures seen from the side of the method rather than the side of the approximation. What the occupancy closure is applied to is set by the recursion axis, and what the interval-signal closure is applied to is set by the window axis. One conditioning choice governs both, and this is worth stating before the algebra rather than after it: the states the interval is conditioned on are the states the interval-averaged conductance is computed from *and* the states the Bayesian update is taken on, so moving along the window axis moves the conductance and the gain together. Both members of that statement are set out at the end of this section, once the equations they differ in are on the page: Table 1 places all eight against those equations, where the axes can be read off as settings rather than as a taxonomy, and Figure 1 draws six of them as one turn of the filter cycle each.

Two things follow from the window axis. First, IR occupies the top rung below the intractable exact likelihood, so its standing in the family is structural and set before any measurement. Second, the lower rungs are established likelihoods rather than constructed foils: the non-recursive instantaneous member is the independent-interval likelihood of *Milescu et al. (2005)*, and the recursive instantaneous member is the macroscopic recursive filter of *Moffatt (2007)* and the generalized filter of *Münch et al. (2022)*. The lattice is one coordinate system that contains these methods as special cases.

The boundary state

A *boundary state* is the pair $(i \rightarrow j)$ of a channel's states at the two ends of an acquisition interval. Chaining intervals, the end state of one interval is the start state of the next, so there is a single state variable per junction: the filter conditions on the states at the interval boundaries and marginalizes the trajectory between them analytically. This is the same device as static condensation, or the spatial Markov property, applied on the time axis. It is not a transition state in the mechanistic sense, and the word *transition* is not used for it here. Because the joint dynamics of the two endpoints are generated by the Kronecker sum $\mathbf{Q} \oplus \mathbf{Q} = \mathbf{Q} \otimes \mathbf{I} + \mathbf{I} \otimes \mathbf{Q}$, whose propagator factorizes as $e^{(\mathbf{Q} \oplus \mathbf{Q})t} = e^{\mathbf{Q}t} \otimes e^{\mathbf{Q}t}$, the moments over the K^2 boundary pairs collapse to bilinear forms in the K -dimensional state space and no explicit K^2 object is formed (Methods).

The ladder needs two conductance objects, both built from the boundary-conditioned single-channel moments $\bar{\gamma}_{i \rightarrow j}$ (the mean interval-averaged conductance of a channel that starts in i and ends in j) and $\bar{v}_{i \rightarrow j}$ (its variance). Both are moments of the integrated conductance $A_{0 \rightarrow \Delta}$ of Eq. 2, taken conditional on the chain starting in i and ending in j , and both are available in closed form, with no object larger than $K \times K$ formed anywhere. They are derivatives at $\alpha = 0$ of the tilted

semigroup met above. Writing

$$\mathbf{F}_1(\Delta) = \left. \frac{\partial}{\partial \alpha} e^{(\mathbf{Q} + \alpha \mathbf{\Gamma})\Delta} \right|_{\alpha=0}, \quad 2\mathbf{F}_2(\Delta) = \left. \frac{\partial^2}{\partial \alpha^2} e^{(\mathbf{Q} + \alpha \mathbf{\Gamma})\Delta} \right|_{\alpha=0}, \quad (3)$$

with $\mathbf{\Gamma} = \text{diag}(\gamma)$ as above, the $(i \rightarrow j)$ entry of the first is $E[A_{0 \rightarrow \Delta} \mathbf{1}\{X_\Delta = j\} \mid X_0 = i]$ and of the second the same expectation with $A_{0 \rightarrow \Delta}^2$ in place of $A_{0 \rightarrow \Delta}$; the factor of two is the two triangular halves of the double integral the second derivative unfolds into. Dividing by the transition probability conditions on the endpoints and dividing by Δ turns the integral into an average, so

$$\bar{\gamma}_{i \rightarrow j} = \frac{[\mathbf{F}_1(\Delta)]_{i \rightarrow j}}{\Delta P_{i \rightarrow j}(\Delta)}, \quad \bar{v}_{i \rightarrow j} = \frac{2[\mathbf{F}_2(\Delta)]_{i \rightarrow j}}{\Delta^2 P_{i \rightarrow j}(\Delta)} - \bar{\gamma}_{i \rightarrow j}^2. \quad (4)$$

Both derivatives are evaluated from an eigendecomposition of the generator, which is the route the implementation takes and the one the P2X2 analysis of *Moffatt and Pierdominici-Sottile (2025)* was run on. Writing $\mathbf{Q} = \mathbf{V}\mathbf{\Lambda}\mathbf{W}$, with $\mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_K)$ the eigenvalues, $\mathbf{V}$ the matrix of right eigenvectors and $\mathbf{W} = \mathbf{V}^{-1}$, the transition matrix is $\mathbf{P}(\Delta) = \mathbf{V}e^{\mathbf{\Lambda}\Delta}\mathbf{W}$ and the two derivatives are the same sandwich with the conductance carried in the eigenbasis. With $\hat{\mathbf{\Gamma}} = \mathbf{W}\mathbf{\Gamma}\mathbf{V}$ and $\circ$ the entrywise product,

$$\frac{\mathbf{F}_1(\Delta)}{\Delta} = \mathbf{V}(\hat{\mathbf{\Gamma}} \circ \mathbf{E}_2)\mathbf{W}, \quad (\mathbf{E}_2)_{ab} = E_2(\lambda_a \Delta, \lambda_b \Delta), \quad (5)$$

$$\frac{2\mathbf{F}_2(\Delta)}{\Delta^2} = 2\mathbf{V}\mathbf{M}\mathbf{W}, \quad M_{ab} = \sum_c \hat{\Gamma}_{ac} \hat{\Gamma}_{cb} E_3(\lambda_a \Delta, \lambda_c \Delta, \lambda_b \Delta), \quad (6)$$

where E_2 and E_3 are the first and second divided differences of the exponential,

$$E_2(x, y) = \frac{e^x - e^y}{x - y}, \quad E_3(x, y, z) = \frac{e^x}{(x - y)(x - z)} + \frac{e^y}{(y - x)(y - z)} + \frac{e^z}{(z - x)(z - y)}. \quad (7)$$

Both are symmetric in their arguments and both are smooth where the denominators vanish, so repeated and near-repeated eigenvalues are read off the coincidence limits $E_2(x, x) = e^x$ and $E_3(x, x, x) = e^x/2$, with the one-pair-coincident form between them (Supplementary Information). Eq. 5 is the form in which the first moment is given in the supplement of *Moffatt and Pierdominici-Sottile (2025)*; the second moment was not stated there. The eigendecomposition is computed once per agonist concentration and reused by every window at that concentration, after which each window costs the $O(K^3)$ contraction of Eq. 6.

An equivalent evaluation avoids diagonalizing $\mathbf{Q}$ altogether. $\mathbf{F}_1$ and $\mathbf{F}_2$ are the strictly upper blocks of the exponentials of the augmented generators

$$\mathcal{Q}_2 = \begin{pmatrix} \mathbf{Q} & \mathbf{\Gamma} \\ \mathbf{0} & \mathbf{Q} \end{pmatrix}, \quad \mathcal{Q}_3 = \begin{pmatrix} \mathbf{Q} & \mathbf{\Gamma} & \mathbf{0} \\ \mathbf{0} & \mathbf{Q} & \mathbf{\Gamma} \\ \mathbf{0} & \mathbf{0} & \mathbf{Q} \end{pmatrix}, \quad (8)$$

of sizes $2K \times 2K$ and $3K \times 3K$, whose exponentials are block triangular with $e^{\mathbf{Q}\Delta}$ on the diagonal (*Van Loan, 1978*), so that $e^{\mathbf{Q}_3\Delta}$ on its own returns the transition matrix and both moments. It identifies the two objects with no spectral machinery, and where no reliable eigensolver is at hand it is the construction to use. The route taken here is the spectral one, and the reason is the derivatives rather than the exponential. This paper propagates the derivative of the log-likelihood with respect to the kinetic parameters through every object the filter builds, and the eigenvalues and eigenvectors of $\mathbf{Q}$ move with those parameters by a first-order perturbation of the same decomposition, in closed form, while a route built on $e^{\mathbf{Q}_3\Delta}$ carries that derivative payload through an object of side $3K$ whose nine blocks hold three distinct matrices. Every number reported in this paper comes from the spectral route.

Both expressions are 0/0 where a transition is unreachable, which is not a numerical edge case here: at zero agonist $k_{\text{on}}[A]$ vanishes identically, so $P_{C \rightarrow O}(\Delta) = 0$ exactly over the whole pre-pulse

segment. Conditioning on a transit that occupies a vanishing fraction of the window leaves the chain at its two endpoint states for essentially equal times, so the limit is the endpoint average,

$$\bar{\gamma}_{i \rightarrow j} \rightarrow \frac{1}{2}(\gamma_i + \gamma_j), \quad \bar{v}_{i \rightarrow j} \rightarrow \left(\frac{\gamma_i - \gamma_j}{2} \right)^2, \quad (9)$$

which depends on the conductance vector alone. The implementation does not branch on the singular case: it reads the two quotients of Eq. 4 as posterior means and shrinks them towards this limit with a pseudo-count whose weight is the floating-point noise of the spectral reconstruction, so the limit is approached smoothly and no fitted constant enters (Methods). The start-conditioned mean conductance is the K -vector

$$\bar{\gamma}_{i \rightarrow \cdot} = \sum_j P_{i \rightarrow j}(\Delta) \bar{\gamma}_{i \rightarrow j}, \quad (10)$$

which conditions on the state at the interval's start alone and is the object `INR`, `MR` and `VR` use. The boundary-conditioned mean conductance is the same average retained on both endpoints, so it stays a $K \times K$ object rather than collapsing to a vector, and is the object `IR` uses.

The interval update

The interval likelihood is the instantaneous one applied to the boundary state. The recursive member of [Moffatt \(2007\)](#) carries a Gaussian over the channel occupancies and conditions it on each recorded sample: in the per-channel normalization used throughout, with γ the per-state conductance and ϵ^2 the instrumental variance,

$$\begin{aligned} y^{\text{pred}} &= N_{\text{ch}} \boldsymbol{\mu}^T \boldsymbol{\gamma}, & \mathbf{g} &= \boldsymbol{\Sigma} \boldsymbol{\gamma}, & \sigma^2 &= \epsilon^2 + N_{\text{ch}} \boldsymbol{\gamma}^T \boldsymbol{\Sigma} \boldsymbol{\gamma}, \\ \boldsymbol{\mu}^{\text{post}} &= \boldsymbol{\mu} + \frac{\mathbf{g}}{\sigma^2} \delta, & \boldsymbol{\Sigma}^{\text{post}} &= \boldsymbol{\Sigma} - \frac{N_{\text{ch}} \mathbf{g} \mathbf{g}^T}{\sigma^2}, & \delta &= y^{\text{obs}} - y^{\text{pred}}. \end{aligned} \quad (11)$$

This is the Gaussian conditioning of the pair (state, observation) on the value recorded. The gain $\mathbf{g}$ is the covariance between the state a channel is in and the current it carries, and it is the whole of what brings the observation to the state: the predicted variance divides it, and nothing else in the update touches the data. It is also the object the interval version has to rebuild, since the two contractions of Eq. 11 against $\boldsymbol{\Sigma}$ are $\mathbf{g}$ itself and $\boldsymbol{\gamma}^T \mathbf{g}$, the same product taken once and twice. Every moment in this section belongs to a single channel. Writing $\mathbf{X}$ for the indicator vector of the state one channel occupies, which is one on that state and zero elsewhere, $\boldsymbol{\mu} = \mathbb{E}[\mathbf{X}]$ is its probability vector and $\boldsymbol{\Sigma} = \text{Cov}(\mathbf{X})$ its covariance, which is $\text{diag}(\boldsymbol{\mu}) - \boldsymbol{\mu} \boldsymbol{\mu}^T$ for a channel nothing has been observed about and vanishes when its state is known. Both are objects of one channel, and neither is a population quantity divided by anything. Channels are independent of each other, so the counts $\mathbf{N} = \sum_c \mathbf{X}^{(c)}$ over the ensemble have $\mathbb{E}[\mathbf{N}] = N_{\text{ch}} \boldsymbol{\mu}$ and $\text{Cov}(\mathbf{N}) = N_{\text{ch}} \boldsymbol{\Sigma}$, and that is where every factor of N_{ch} in Eq. 11 comes from: the recorded current is a population quantity and the state it is informative about is carried per channel.

One generalization, and the whole family is inside it.

Eq. 11 takes a channel's conductance to be a number once its state is known. That is true of the current at an instant and of nothing else. One step of generality covers every member below: alongside the conductance, supply a *per-state variance* $\mathbf{w}$, the variance the recorded value still has once the state is fixed. Channels being independent given their states, those variances add over the ensemble, so

$$\mathbb{E}[\bar{y} \mid \mathbf{N}] = \boldsymbol{\gamma}^T \mathbf{N}, \quad \text{Var}[\bar{y} \mid \mathbf{N}] = \epsilon_{0 \rightarrow \Delta}^2 + \sum_s N_s w_s, \quad (12)$$

and the extra variance is linear in the counts, exactly as the mean is. Taking the expectation over the state distribution turns the predictive variance of Eq. 11 into

$$\sigma^2 = \epsilon_{0 \rightarrow \Delta}^2 + N_{\text{ch}} \boldsymbol{\gamma}^T \boldsymbol{\Sigma} \boldsymbol{\gamma} + N_{\text{ch}} \boldsymbol{\mu}^T \mathbf{w}, \quad (13)$$

with every other line of Eq. 11 unchanged. Setting $\mathbf{w} = \mathbf{0}$ returns the published filter. This is the same device *Münch et al. (2022)* add to the same filter to model open-channel noise, from a different motivation; here the state-dependent term will carry the gating variance the chosen state leaves unexplained rather than a property of the pore.

Four contractions, and the members differ in three of them.

Written this way, Eq. 11 contracts a member's objects against its own state distribution and covariance, and it does so four times:

$$\underbrace{\boldsymbol{\mu}^\top \boldsymbol{\gamma}}_{\text{the mean}} \quad \underbrace{\boldsymbol{\mu}^\top \mathbf{w}}_{\text{the state-dependent noise}} \quad \underbrace{\boldsymbol{\gamma}^\top \boldsymbol{\Sigma} \boldsymbol{\gamma}}_{\text{what the state contributes}} \quad \underbrace{\boldsymbol{\Sigma} \boldsymbol{\gamma}}_{\text{state against current, } \mathbf{g}}$$

two against $\boldsymbol{\mu}$ and two against $\boldsymbol{\Sigma}$, and everything else in Eq. 11 is a scalar. Extending the filter to a recorded value that spans the window therefore cannot require a new algorithm. It requires choosing a state variable and supplying $(\boldsymbol{\gamma}, \mathbf{w})$ on it, and the members below are those choices. The per-state variance is not a free third choice beside the other two: Eq. 12 defines it as whatever variance the chosen state leaves unexplained, so naming the state names it.

What Eq. 13 gets exactly right, and what it gives up.

The three moments the update consumes are exact functions of $(\boldsymbol{\mu}, \boldsymbol{\Sigma})$. In particular the per-state variance enters σ^2 and touches neither the predicted mean nor $\mathbf{g}$: given the counts, the extra fluctuation has mean zero, so $\text{Cov}(\mathbf{N}, \bar{y}) = \text{Cov}(\mathbf{N}, \mathbb{E}[\bar{y} \mid \mathbf{N}])$ and $\mathbf{w}$ drops out of it. That is what makes charging the leftover variance to the noise term legitimate rather than convenient, and it is used at Eq. 19 below.

What is given up is not exactness but *sufficiency*. Because $\text{Var}[\bar{y} \mid \mathbf{N}]$ depends on $\mathbf{N}$, the prior over the state enters the variance of the likelihood and not only its mean, and a normal prior is conjugate for the mean of a normal likelihood but not for its variance. The pair $(\mathbf{N}, \bar{y})$ is therefore not jointly Gaussian, and the update of Eq. 11 is not the exact Bayesian update: it conditions a Gaussian that carries the exact first two moments, treating a heteroscedastic emission as if it were homoscedastic. *Münch et al. (2022)* meet the same obstruction and resolve it the same way, noting that the state-dependent term makes the incremental likelihood intractable because “the conjugacy property is lost”, applying the total-variance decomposition to recover mean and covariance, and stating the cost in the terms used here: the covariance equations “are still exact for the more general noise model”, while “mean and covariance just do not form a set of sufficient statistics anymore”. Two groups reached the same device from different motivations, which is worth more to a reader than either group's assurance.

What stops Eq. 11 from being used on the interval average is that the average is not a function of the occupancy at any one time, so no conductance vector on the K states delivers its mean. On the K^2 boundary pairs one does: a channel that starts in i and ends in j contributes $\bar{y}_{i \rightarrow j}$ on average, so the interval average is linear in the boundary counts again. That is the whole of the move, and three steps finish it. Run Eq. 11 on the pairs, with $\bar{y}_{i \rightarrow j}$ in place of $\boldsymbol{\gamma}$ and the boundary covariance in place of $\boldsymbol{\Sigma}$. Charge what the conductance still varies by once the pair is fixed to ϵ^2 , as though the recording carried an extra noise of variance $\bar{v}_{i \rightarrow j}$ per channel, of a size that depends on the pair that channel took. Marginalize over the start index, which the next window has no use for. Every equation in this subsection is one of those three steps carried out, and the two objects the interval likelihood is usually written with are those two contractions read on the larger space.

The boundary state's own Gaussian follows from the prior and the propagator. Let $\mathbf{X}_{0 \rightarrow \Delta}$ be the indicator of the pair a single channel takes, one on the (i, j) it starts and ends in and zero on the other $K^2 - 1$, and let $(\boldsymbol{\mu}_0, \boldsymbol{\Sigma}_0)$ be the channel's state at the window's start. It leaves i and lands in j with probability $P_{i \rightarrow j}(\Delta)$, so

$$\boldsymbol{\mu}_{i \rightarrow j} = \boldsymbol{\mu}_{0,i} P_{i \rightarrow j}(\Delta), \quad (14)$$

$$(\boldsymbol{\Sigma}_{0 \rightarrow \Delta})_{(i \rightarrow j), (k \rightarrow l)} = P_{i \rightarrow j}(\Delta) (\boldsymbol{\Sigma}_0 - \text{diag}(\boldsymbol{\mu}_0))_{ik} P_{k \rightarrow l}(\Delta) + [\text{diag}(\boldsymbol{\mu}_{0 \rightarrow \Delta})]_{(i,j), (k,l)}. \quad (15)$$

The two terms of Eq. 15 are the two sources of spread. The first carries the uncertainty already in the prior across the window, once on each index. The second is the uncertainty a channel known to be leaving i still has about where it lands, the covariance of a single draw from row i of $\mathbf{P}(\Delta)$, which is diagonal on the pairs because a channel takes exactly one of them. Subtracting $\text{diag}(\boldsymbol{\mu}_0)$ in the first term is what keeps the prior's own single-channel part from being counted twice, since the second term regenerates it at the resolution of the pairs.

Summing the start index out of Eqs. 14 and 15 gives the occupancy Gaussian at the end of the window,

$$\boldsymbol{\mu}_\Delta = \mathbf{P}(\Delta)^\top \boldsymbol{\mu}_0, \quad (16)$$

$$\boldsymbol{\Sigma}_\Delta = \mathbf{P}(\Delta)^\top (\boldsymbol{\Sigma}_0 - \text{diag}(\boldsymbol{\mu}_0)) \mathbf{P}(\Delta) + \text{diag}(\boldsymbol{\mu}_\Delta), \quad (17)$$

term by term, the diagonal of the one becoming the diagonal of the other. This is the propagation an unconditioned member performs, and it is Eq. 15 seen from the end of the window alone. The first line of Eq. 11, on the boundary state, is the predictive. Its mean is

$$\bar{y}_{0 \rightarrow \Delta}^{\text{pred}} = N_{\text{ch}} \boldsymbol{\mu}_0^\top \bar{\mathbf{y}}_0, \quad (18)$$

since $\boldsymbol{\mu}_{0 \rightarrow \Delta}$ contracted against $\bar{\mathbf{y}}$ is Eq. 10 summed against the prior, and its variance has three terms,

$$\sigma_{0 \rightarrow \Delta}^2 = \epsilon_{0 \rightarrow \Delta}^2 + N_{\text{ch}} \widetilde{\boldsymbol{\gamma}^\top \boldsymbol{\Sigma} \boldsymbol{\gamma}} + N_{\text{ch}} \boldsymbol{\mu}_0^\top \mathbf{w}. \quad (19)$$

The first is the ϵ^2 of Eq. 11, the instrumental noise now taken over the window, which carries the acquisition averaging and decreases with Δ ; its parametric form is given in Methods. The third is the noise charged in step two, $N_{\text{ch}} \sum_{i,j} \boldsymbol{\mu}_{i \rightarrow j} \bar{v}_{i \rightarrow j}$, which is what Eq. 19 writes with $\mathbf{w}$, the start-conditioned interval variance, once the end index is summed. The middle term is $\boldsymbol{\gamma}^\top \boldsymbol{\Sigma} \boldsymbol{\gamma}$ of Eq. 11 evaluated on the pairs, and the tilde is what marks that: under it the two $\boldsymbol{\gamma}$ stand for $\bar{\mathbf{y}}_{i \rightarrow j}$ and $\boldsymbol{\Sigma}$ for the $K^2 \times K^2$ boundary covariance, not for the K -vector and the $K \times K$ matrix they name everywhere else. Written out with $\boldsymbol{\Sigma}_{0 \rightarrow \Delta}$ of Eq. 15 and then collapsed by substituting it,

$$\begin{aligned} \widetilde{\boldsymbol{\gamma}^\top \boldsymbol{\Sigma} \boldsymbol{\gamma}} &= \sum_{i,j} \sum_{k,l} \bar{\mathbf{y}}_{i \rightarrow j} (\boldsymbol{\Sigma}_{0 \rightarrow \Delta})_{(i,j),(k,l)} \bar{\mathbf{y}}_{k \rightarrow l} \\ &= \bar{\mathbf{y}}_0^\top (\boldsymbol{\Sigma}_0 - \text{diag}(\boldsymbol{\mu}_0)) \bar{\mathbf{y}}_0 + \boldsymbol{\mu}_0^\top (\mathbf{H} \mathbf{1}), \end{aligned} \quad (20)$$

where $H_{i \rightarrow j} = \bar{\mathbf{y}}_{i \rightarrow j}^\top P_{i \rightarrow j}(\Delta)$ and $\mathbf{1}$ is the vector of ones. The two summands are the two terms of Eq. 15 in order, and only K -vectors and $K \times K$ matrices survive, so the $K^2 \times K^2$ covariance is never built. Reading the tilde as $\bar{\mathbf{y}}_0^\top \boldsymbol{\Sigma}_0 \bar{\mathbf{y}}_0$ keeps the first summand and drops the second, which is the single-channel spread of the interval mean across boundary pairs and does not vanish when the occupancy covariance does.

The gain is the same contraction taken once. Built on the pairs it is the K^2 -vector $\sum_{k,l} (\boldsymbol{\Sigma}_{0 \rightarrow \Delta})_{(i,j),(k,l)} \bar{\mathbf{y}}_{k \rightarrow l}$, and step three sums its start index away, leaving the $\mathbf{g}_{0 \rightarrow \Delta}$ the interval update uses,

$$\mathbf{g}_{0 \rightarrow \Delta} = \widetilde{\boldsymbol{\gamma}^\top \boldsymbol{\Sigma}} = \mathbf{P}(\Delta)^\top (\boldsymbol{\Sigma}_0 - \text{diag}(\boldsymbol{\mu}_0)) \bar{\mathbf{y}}_0 + \mathbf{G}^\top \boldsymbol{\mu}_0, \quad \mathbf{G}_{i \rightarrow j} = \bar{\mathbf{y}}_{i \rightarrow j}^\top P_{i \rightarrow j}(\Delta), \quad (21)$$

the covariance between the state $\mathbf{X}_\Delta$ a channel ends the window in and the current that same channel contributes to the window average. Over the ensemble it is $\text{Cov}(\mathbf{N}_\Delta, \bar{\mathbf{y}}) = N_{\text{ch}} \mathbf{g}_{0 \rightarrow \Delta}$, with $\mathbf{N}_\Delta$ the counts by state at the end of the window, and that is the form the update uses. Again the two summands are the two terms of Eq. 15: $\mathbf{G}$ carries one power of $\bar{\mathbf{y}}$ where $\mathbf{H}$ carried two, and the end index is left open instead of being summed against $\mathbf{1}$.

The second line of Eq. 11, marginalized the same way, is the update, with $\delta = \bar{y}_{0 \rightarrow \Delta}^{\text{obs}} - \bar{y}_{0 \rightarrow \Delta}^{\text{pred}}$ the innovation,

$$\boldsymbol{\mu}^{\text{post}} = \boldsymbol{\mu}_\Delta + \frac{\mathbf{g}_{0 \rightarrow \Delta}}{\sigma_{0 \rightarrow \Delta}^2} \delta, \quad (22)$$

$$\boldsymbol{\Sigma}^{\text{post}} = \boldsymbol{\Sigma}_\Delta - \frac{N_{\text{ch}} \mathbf{g}_{0 \rightarrow \Delta} \mathbf{g}_{0 \rightarrow \Delta}^\top}{\sigma_{0 \rightarrow \Delta}^2}. \quad (23)$$

Conditioning and summing the start index commute here, because the sum is linear and the correction is rank one on the index it acts on, so the downdate crosses it unchanged, with $\mathbf{g}_{0 \rightarrow \Delta}$ standing where the boundary gain stood. What the sum discards is the posterior correlation between the interval current and the state the window *started* in, which costs nothing inside the window and is what the recursion pays for below.

Two of the three steps are exact and one is the approximation. Eqs. 14 and 15 are the true moments of $\mathbf{X}_{0 \rightarrow \Delta}$, and so are the three the predictive consumes. Writing $\mathbf{N}_{0 \rightarrow \Delta} = \sum_c \mathbf{X}_{0 \rightarrow \Delta}^{(c)}$ for the boundary counts over the ensemble, channels are independent given the pairs each of them took, so given $\mathbf{N}_{0 \rightarrow \Delta}$ the recorded average has mean $\sum_{i,j} \bar{y}_{i \rightarrow j} N_{i \rightarrow j}$ and variance $\epsilon_{0 \rightarrow \Delta}^2 + \sum_{i,j} \bar{v}_{i \rightarrow j} N_{i \rightarrow j}$ exactly, and the state-dependent noise of step two, having no covariance with the counts that carry it, enters $\sigma_{0 \rightarrow \Delta}^2$ and changes neither $\bar{y}_{0 \rightarrow \Delta}^{\text{pred}}$ nor $\mathbf{g}_{0 \rightarrow \Delta}$. That last point is what makes charging it to the noise term legitimate rather than convenient. What is approximate is the Gaussian, in the two places the box above names: the counts are not Gaussian, which is the occupancy closure, and neither is the window average of one channel given its pair, so replacing its law by the two numbers $\bar{y}_{i \rightarrow j}$ and $\bar{v}_{i \rightarrow j}$ is the interval-signal closure. The two together still do not make the pair $(\mathbf{N}_{0 \rightarrow \Delta}, \bar{y})$ jointly Gaussian, since its conditional variance depends on the state, so the update conditions the Gaussian that carries the exact moments. The factor N_{ch} in the downdate, and its absence from the mean correction, come from the same place. All three of $\boldsymbol{\mu}$, $\boldsymbol{\Sigma}$ and $\mathbf{g}_{0 \rightarrow \Delta}$ belong to one channel, while what is observed belongs to the ensemble, so the conditioning is done on $\text{Cov}(\mathbf{N}_{\Delta}, \bar{y}) = N_{\text{ch}} \mathbf{g}_{0 \rightarrow \Delta}$ against $\text{Cov}(\mathbf{N}_{\Delta}) = N_{\text{ch}} \boldsymbol{\Sigma}_{\Delta}$. The rank-one term picks up the channel count twice and the return to one channel gives one back, leaving a single factor; the mean correction picks it up once and gives it back, leaving none. In a recursive member the posterior pair $(\boldsymbol{\mu}^{\text{post}}, \boldsymbol{\Sigma}^{\text{post}})$ becomes the prior for the next window; in a non-recursive member it is not carried forward, and each window's predictive is propagated from the initial condition alone. Reusing the posterior as the next prior is itself a closure: it is exact only while the end marginal is a sufficient summary of the boundary posterior, which the Gaussian projection does not guarantee. For this reason the recursive members' Gaussian information is a model-based quantity and is not claimed to be the exact information of the process; how far it departs from the truth is what the Diagnostics measure.

The start-conditioned members, stated forward

The construction above reaches IR by choosing the boundary pair as the state. The members that stop at the interval's start are usually presented as that construction with the collapse brought forward, and read that way they look like abbreviations of it. They are not, and the difference matters for what is being compared later in this section, so each is stated here from its own assumption instead.

MR takes the recorded value to depend on the occupancy where the window opens. Formally this asserts

$$\bar{y}_{0 \rightarrow \Delta} \perp X_{\Delta} \mid X_0, \quad (24)$$

that once the state a channel started in is known, the recorded average says nothing further about where that channel ended. Everything MR does follows from Eq. 24, and the member is complete without any object larger than the K states. What it must supply is a conductance and a per-state variance on that state, and both are conditional moments of $A_{0 \rightarrow \Delta}$ given the start alone:

$$(\bar{y}_0)_i = \mathbb{E} \left[\frac{1}{\Delta} A_{0 \rightarrow \Delta} \mid X_0 = i \right], \quad (\sigma_{\bar{y}_0}^2)_i = \text{Var} \left[\frac{1}{\Delta} A_{0 \rightarrow \Delta} \mid X_0 = i \right], \quad (25)$$

the first being Eq. 10 and the second the per-state variance Eq. 13 then asks for. Neither needs a state larger than K . Both come from the same tilted semigroup of Eq. 3 that any interval member needs, but from its derivatives contracted against the vector of ones rather than read entry by entry: what MR requires are the row sums, and the row sums are integrals of K -vectors. This is worth stating plainly because it is what makes MR a construction in its own right, reached without ever forming a pair.

Eq. 24 is an assumption about the process and can be examined on its own terms. It is false for any window of nonzero length: a channel that starts shut and delivers current during the window is more likely to have ended open than one that delivers none, so the recorded average does carry information about the end state beyond what the start state gives. IR is the member that declines to assume it, and the boundary pair is the state on which the recorded average is linear with no conditional expectation taken at all. MR is therefore a correctly derived filter for a measurement model that is false, rather than an approximate version of IR ; how much the falsehood costs is Eq. 30 below.

What separates the start-conditioned members from the boundary one

Run as the same three steps on the K states instead of the pairs, with $\bar{\gamma}_0$ in place of γ , two algebraic differences separate them from the one that conditions on both ends. The first is the form of the third term of Eq. 19. MR carries the total per-start-state interval variance, the spread of the interval-averaged conductance of a channel known only to have started in i , while VR and IR carry the residual variance left once the end state is also conditioned on, $\sum_j P_{i \rightarrow j}(\Delta) \bar{v}_{i \rightarrow j}$. By the law of total variance the two forms differ by the spread of the interval mean across end states,

$$\underbrace{\sum_j P_{i \rightarrow j}(\Delta) [\bar{v}_{i \rightarrow j} + (\bar{\gamma}_{i \rightarrow j} - \bar{\gamma}_{i \rightarrow \cdot})^2]}_{\text{total, MR}} = \underbrace{\sum_j P_{i \rightarrow j}(\Delta) \bar{v}_{i \rightarrow j}}_{\text{residual, VR, IR}} + \text{Var}_j [\bar{\gamma}_{i \rightarrow j} | i], \quad (26)$$

so MR charges to the observation variance a spread that the end state, once conditioned on, would resolve. The second difference is the boundary cross-covariance term

$$N_{\text{ch}} \gamma^T \Sigma \gamma \quad (27)$$

which IR retains and both start-conditioned members drop. It enters twice, in the second term of Eq. 19 and in the gain $\mathbf{g}_{0 \rightarrow \Delta}$ of Eqs. 22 and 23.

The two differences are not independent, and the tempting reading of them, one step that moves the variance followed by one that moves the gain, is wrong in both directions. Taken from a common prior they cancel in the predicted variance: what the boundary term adds to the second term of Eq. 19 is $\mu_0^T \text{Var}_j [\bar{\gamma}_{i \rightarrow j} | i]$, which is what the residual form removes from the third, so from the same state MR and IR assign the interval average the same variance by two different decompositions and the whole of what separates them is the gain (Appendix 1). In the other direction, the predicted variance divides the gain, so changing the variance form changes the update as well: VR removes that spread from the third term without restoring it in the second, predicts less variance than either from the same state, and its posterior departs from MR 's at once. The equality is therefore between two maps from a prior to a predicted variance, and not between two recordings. Once the gains differ the priors differ, and the variances the two members report along one recording differ with them, which is what the Results measure across the regime grid.

Both members against the exact law

Neither member is measured against the other. Both are approximations to one joint law, that of $(\mathbf{N}_\Delta, \bar{y})$ under the true process, and given (μ_0, Σ_0) that law has exact first two moments which either member can be held against. Two elementary decompositions supply them, and they say different things about the two blocks.

The variance: both members are exact, and their slots differ.

By the law of total variance, conditioning on the counts where the window opens,

$$\text{Var}[\bar{y}] = \epsilon_{0 \rightarrow \Delta}^2 + \underbrace{\mathbb{E}[\text{Var}(\bar{y} | \mathbf{N}_0)]}_{N_{\text{ch}} \mu_0^T \mathbf{w}} + \underbrace{\text{Var}[\mathbb{E}(\bar{y} | \mathbf{N}_0)]}_{N_{\text{ch}} \bar{\gamma}_0^T \Sigma_0 \bar{\gamma}_0}, \quad (28)$$

which is Eq. 19 read with the start-conditioned objects: the middle term is $\tilde{\mathbf{y}}_0^\top \Sigma_0 \tilde{\mathbf{y}}_0$, the tilde dropped, and $\mathbf{w}$ is the total slot of Eq. 26. Read with the boundary objects instead, the tilde of Eq. 20 in the middle and the residual slot in $\mathbf{w}$, the same quantity is Eq. 19 again. **Both members report the exact variance of the recorded value** given (μ_0, Σ_0) , so they agree with each other for the reason that each agrees with the truth, and nothing is measured by that agreement. This is Eq. 26 read from the outside: the same second moment sits in the state term of one member and in the noise term of the other, which is why the totals coincide.

The variable conditioned on in Eq. 28 is the count vector and not one channel's state. Eq. 12 gives $\text{Var}[\bar{y} | \mathbf{N}_0] = \epsilon_{0 \rightarrow \Delta}^2 + \sum_s N_{0,s} w_s$ and $\mathbb{E}[\bar{y} | \mathbf{N}_0] = \tilde{\mathbf{y}}_0^\top \mathbf{N}_0$, so the second underbrace is $\tilde{\mathbf{y}}_0^\top \text{Cov}(\mathbf{N}_0) \tilde{\mathbf{y}}_0$ and the per-channel convention supplies the factor of N_{ch} in both. Conditioning on a single channel's X_0 returns the same expressions only while Σ_0 is still that channel's own covariance, which it stops being at the first update.

The covariance: one member is exact and the other is short by a named amount. By the law of total covariance, conditioning on the same counts and carried on the per-channel scale,

$$\frac{1}{N_{\text{ch}}} \text{Cov}[\mathbf{N}_\Delta, \bar{y}] = \underbrace{\frac{1}{N_{\text{ch}}} \text{Cov}[\mathbb{E}(\mathbf{N}_\Delta | \mathbf{N}_0), \bar{y}]}_{\mathbf{g}_{0 \rightarrow \Delta}^{\text{MR}}} + \frac{1}{N_{\text{ch}}} \mathbb{E}[\text{Cov}(\mathbf{N}_\Delta, \bar{y} | \mathbf{N}_0)]. \quad (29)$$

The first term is the start-conditioned gain $\mathbf{g}_{0 \rightarrow \Delta}^{\text{MR}} = \mathbf{P}(\Delta)^\top \Sigma_0 \tilde{\mathbf{y}}_0$ exactly. The second is what Eq. 24 sets to zero: given the counts the channels are conditionally independent, so it is the per-channel conditional covariance summed over the population, and a conditional independence is precisely the statement that it vanishes. It does not vanish. IR carries the whole of Eq. 29, so the difference between the two members in this block is that second summand and nothing else,

$$\mathbf{g}_{0 \rightarrow \Delta}^{\text{IR}} - \mathbf{g}_{0 \rightarrow \Delta}^{\text{MR}} = \left(\sum_i \mu_{i \rightarrow j} D_{i \rightarrow j} \right)_j, \quad D_{i \rightarrow j} = \bar{y}_{i \rightarrow j} - (\tilde{y}_0)_i, \quad (30)$$

which is Eq. 21 minus $\mathbf{P}(\Delta)^\top \Sigma_0 \tilde{\mathbf{y}}_0$ written out. The deviation $D_{i \rightarrow j}$ is a property of the process and not of either algorithm: it is how far the interval average of a channel that ended in j sits from the average over all the places it might have ended. It has zero row sums, $\sum_j P_{i \rightarrow j}(\Delta) D_{i \rightarrow j} = 0$, since $(\tilde{y}_0)_i$ is that very average, and it therefore has zero mean under $\mu_{i \rightarrow j}$ as well, so the gap between the two forms of Eq. 26 is exactly its variance under the law of the pairs. Eq. 30 does not depend on Σ_0 at all.

Why the same quantity is invisible in one block and decisive in the other.

D enters the variance squared and the covariance once, and the two blocks answer to its zero row sums differently. A second moment can sit in either of the two variance terms without the total noticing, which is Eq. 26. A *first* moment cannot: a noise term has no first moment, so when D is collapsed against the end index there is nowhere for it to go, and its column profile, which is what carries the information about where the channel ended, is lost. That asymmetry between D^2 and D is the whole of the difference between the two members, and it is a consequence of Eq. 24 rather than a defect of the arithmetic that implements it.

The same construction in the linear-filtering frame

Everything above was Gaussian conditioning: a joint law over (state, recorded value), its three blocks, and Bayes. Nothing was borrowed from linear filtering, and the derivation does not need it. The correspondence with that literature is nevertheless exact, it is what places this construction against forty years of work on the same measurement problem, and a reader arriving from there should be able to check the identification rather than take it on faith. The relation is conceded in the Discussion; the equations that make it checkable are here.

For a state $\mathbf{x}$ and a scalar observation y jointly Gaussian, the filter's update is

$$\mathbf{x}^{\text{post}} = \mathbf{x}^{\text{prior}} + \frac{\text{Cov}(\mathbf{x}, y)}{\text{Var}(y)} (y - \mathbb{E}[y]), \quad \mathbf{P}^{\text{post}} = \mathbf{P}^{\text{prior}} - \frac{\text{Cov}(\mathbf{x}, y) \text{Cov}(\mathbf{x}, y)^\top}{\text{Var}(y)}, \quad (31)$$

Member	The states are conditioned on	Predicted mean (Eq. 18)	Predicted variance beyond the instrumental term $\epsilon_{0 \rightarrow \Delta}^2$ (Eq. 19)
<i>Open loop</i> : no update; every window is propagated from the initial condition (Eqs. 16, 17 alone)			
LSE	nothing: there are no states	the deterministic mean current at the middle of the window	none: one constant over the whole record, fitted as the residual sum of squares
ILSE	nothing: there are no states	the same mean current averaged over the window	the same constant
NR	an instant, the middle of the window	the current at that instant	the spread of the instantaneous conductance across the states the channels occupy
INR	the state at the start of the window	the interval-mean current given that state	the spread of that interval mean across the occupied states, and the variance of one channel's interval mean given its start state
<i>Recursive</i> : the posterior of one window is the prior of the next (Eqs. 22, 23)			
R	an instant, the middle of the window	the current at that instant	the spread of the instantaneous conductance across the states the channels occupy
MR	the state at the start of the window	the interval-mean current given that state	the spread of that interval mean across the occupied states, and the variance of one channel's interval mean given its start state
VR	the start, and both ends for the interval variance alone	the interval-mean current given the start state	the spread of that interval mean across the occupied states, and the variance of one channel's interval mean given both its endpoints
IR	the states at both ends of the window	the interval-mean current, the unobserved end state marginalized	the spread of the interval mean across the occupied pairs of endpoint states, and the variance of one channel's interval mean given both its endpoints

Table 1. The eight members as settings of the equations above. One conditioning choice runs through each row: the states an interval is conditioned on are the states its conductance is computed from and the states its update is taken on, so the second column governs the gain of Eq. 22 as well as the two columns beside it. Reading down, the second column is the window axis and the two blocks are the recursion axis. Three consequences are worth reading off. LSE and ILSE differ in the mean alone, which isolates the acquisition average with nothing else moving; NR and R, and INR and MR, carry identical Gaussians and differ only in whether the recording is fed back into the states; and MR and IR, taken from the same state, assign the same variance by two decompositions that rearrange into each other, so what separates them is the gain, which is then what makes the variances they report along a recording differ after all, while VR drops the between-endpoint spread from the third term without restoring it in the second and predicts less than either from any state. Read against Eq. 12, that last row is not a fourth reading of the construction but a mismatch: naming the state fixes the per-state variance, and VR fills the slot with the leftover of a state it does not otherwise use, which is why it is the one member whose variance form is a choice rather than a consequence. In print, IR is MacroIR and INR is MacroINR. All eight are implemented and measured: six of them are drawn as one turn of the filter cycle in Figure 1, seven are run over the design plane and the un-averaged least-squares arm at a single cell of it, and the dispatched flags and data keys are in Table 2.




Figure 1. One filter step along the cost ladder. The current from a patch of $N_{\text{ch}} = 20$ two-state channels, recorded as an interval average with white noise, is scored by six members (columns). They form three pairs, one per level of the gating description: least squares on the deterministic mean current (LSE), the open-loop gating likelihood (NR) and the recursive filter (R), each beside its interval-averaged partner, named with an I prefix (ILSE, INR, IR). Rows follow one turn of the cycle: (A) the prior open probability from the Markov step, (B) the predicted current with its spread and the innovation, (C) the Bayesian posterior, absent from the four open-loop columns, and (D) the log-likelihood accumulated from zero. Row (B) is where the members of a pair part company: the left-hand one predicts an instant, a point at the sample with an error bar, the right-hand one the interval average, a segment across the interval with a band. In (A) and (C), IR alone carries a pair of occupancies, one at each end of the interval, tied by a dashed arc: the arc is their covariance and implies no path through the interval. The two-tone disc is the instant that closes one interval and opens the next, and the Markov arrow leaves it because that marginal is all that crosses the boundary. Colour encodes the filter phase (green the channel current, orange predict, blue observe, purple update, black the score), arrows are operations, and a state carried in from the neighbouring interval is faded. The two partial interval corrections, MR and VR, sit off this pairing and are not drawn (Appendix 1). No number is quoted from this 12 ms recording.

Figure 1—figure supplement 1. The same filter cycle over the first ten milliseconds of the recording.

which is Eq. 11 up to the per-channel convention, the source of the factors of N_{ch} there and not here. The two frames do not differ in the update at all: both are the conditioning of a joint Gaussian, and what is usually called the gain is the ratio of the cross block to the observation variance. The entire content of either derivation is the computation of the three blocks $\mathbb{E}[y]$, $\text{Var}(y)$ and $\text{Cov}(\mathbf{x}, y)$ for a recorded value that is an average over the window.

The standard difficulty is that $\bar{y}$ is not a function of the state at any single time, so no observation matrix $\mathbf{H}$ satisfies $\bar{y} = \mathbf{H}\mathbf{x}_\Delta + v$. The device is state augmentation: carry the running integral $z_t = \int_0^t \gamma^\top \mathbf{N}(s) ds$ alongside the counts, so that $(\mathbf{N}, z)$ propagates jointly and the recorded value is z_Δ/Δ plus instrumental noise, which is linear in the augmented state. This is the treatment of an observation that is a flow rather than a stock, and it is a known construction rather than a new one (**Zadrozny, 1988**). Augmentation moves the problem rather than solving it: the filter now needs the variance of the accumulator and its covariance with the occupancy at the end of the window, and in general those are not free. For a continuous-time Markov chain they are, and by the two derivatives already written as Eq. 3: the $(i \rightarrow j)$ entry of $\mathbf{F}_1$ is $\mathbb{E}[A_{0 \rightarrow \Delta} \mathbf{1}\{X_\Delta = j\} \mid X_0 = i]$ and of $2\mathbf{F}_2$ the same with $A_{0 \rightarrow \Delta}^2$, which is exactly what the augmented filter asks for.

Condition on the counts where the window opens; given those the channels are independent, and the three blocks are

$$\mathbb{E}[\bar{y}] = N_{\text{ch}} \boldsymbol{\mu}_0^\top \bar{\mathbf{y}}_0, \quad (32)$$

$$\text{Var}[\bar{y}] = \epsilon_{0 \rightarrow \Delta}^2 + \frac{\text{Var}(z_\Delta)}{\Delta^2}, \quad (33)$$

$$\frac{\text{Cov}[\mathbf{N}_\Delta, \bar{y}]}{N_{\text{ch}}} = \frac{\text{Cov}[\mathbf{N}_\Delta, z_\Delta]}{N_{\text{ch}} \Delta} = \mathbf{P}(\Delta)^\top \boldsymbol{\Sigma}_0 \bar{\mathbf{y}}_0 + \left(\sum_i \mu_{i \rightarrow j} D_{i \rightarrow j} \right)_j. \quad (34)$$

Eq. 32 is Eq. 18, Eq. 33 is Eq. 19, and Eq. 34 is $\mathbf{g}_{0 \rightarrow \Delta}^{\text{MR}}$ plus Eq. 30, which is to say the $\mathbf{g}_{0 \rightarrow \Delta}$ of Eq. 21. The second term of Eq. 34 needs nothing from the boundary construction: for one channel, $\mathbf{F}_1$ divided by Δ and by the transition probability gives $\mathbb{E}[\mathbf{1}\{X_\Delta = j\} \frac{1}{\Delta} A_{0 \rightarrow \Delta} \mid X_0 = i] = P_{i \rightarrow j}(\Delta) \bar{y}_{i \rightarrow j}$, the product of the two conditional expectations is $P_{i \rightarrow j}(\Delta)(\bar{y}_0)_i$, and their difference is $P_{i \rightarrow j}(\Delta) D_{i \rightarrow j}$. The deviation of Eq. 30 is therefore what the first derivative of the tilted semigroup leaves once the product of the two first moments is taken out, and the agreement between the two routes is a result rather than an input.

The augmented-state filter and the boundary-state filter compute the same three blocks, so they perform the same update and return the same posterior over the occupancy. Three things about that follow, and they are the reason the derivation above is given in the other order.

The augmented state has $K + 1$ entries and the boundary state has K^2 , and they deliver the same three blocks. *The pair is a device of the derivation and not information the filter must carry*, which is the same conclusion Eqs. 16 and 17 reached from the other side. The per-state variance disappears in the augmented frame: given $(\mathbf{N}_\Delta, z_\Delta)$ the recorded value is deterministic, so what Eq. 19 splits between its second and third terms sits entirely inside $\text{Var}(z_\Delta)$, and that is why the cross block needs only $\mathbf{F}_1$ while the variance block needs $\mathbf{F}_2$ as well. And the equivalence holds per window, with a reset: the augmented filter must discard the posterior on z and open the next window with $z = 0$, because the next window's integral is a fresh one and is independent of the old given the occupancy.

MR is not the augmented filter under any choice of blocks. It is the filter for a misspecified measurement model,

$$\bar{y} = \mathbf{H}\mathbf{N}_0 + v(\mathbf{N}_0), \quad \mathbf{H} = \bar{\mathbf{y}}_0^\top, \quad \text{Var}[v \mid \mathbf{N}_0] = \epsilon_{0 \rightarrow \Delta}^2 + \mathbf{N}_0^\top \mathbf{w}, \quad (35)$$

whose expectation over $\mathbf{N}_0$ returns Eq. 19. The measurement noise depends on the state, which is outside the textbook filter: with that extension the two frames give the same equations for MR as well, and without it MR has no linear-filtering form at all.

Where the frames genuinely differ is the model rather than the arithmetic. A Kalman derivation presumes a linear-Gaussian state process, and the occupancy counts of a channel population are

not one: they are multinomial, their transition noise depends on the state they leave, and no linear-Gaussian system generates them. The derivation above presumes none of that. It computes the exact first two moments of the true process, over the pairs where the recorded average is genuinely linear, and replaces the law by a Gaussian at the two places the box above names. The two routes arrive at the same equations because the conditioning of a joint Gaussian consumes nothing but the first two moments, and both supply the same ones. Two consequences follow. Naming the update a Kalman correction imports a measurement equation $\bar{y} = \mathbf{H}\mathbf{x} + v$ that the construction never writes down and that is not available for an interval average without the augmentation; and Σ has a meaning in one frame and a convention in the other, since what this section carries is the covariance of one channel's state indicator and not the covariance matrix of a linear system. The pair survives as the presentation even though the augmented state reaches the same blocks with fewer entries, because on the pair the recorded average is linear in the state with no conditional expectation taken, so exactness is structural rather than a separate calculation; the collapse hands over the next window's prior at no cost; and the per-state variance slot is what makes $\mathbf{R}$, $\mathbf{MR}$ and $\mathbf{IR}$ one construction with three readings. The augmentation gives no comparably natural form to the other two members.

Every symbol against the name it carries in the implementation

The equations above are written with the objects the implementation forms, and the correspondence is one to one. It is given because a specification whose symbols cannot be found in the code is not a specification. $\circ$ is the entrywise product and $\mathbf{1}$ the vector of ones; $\bar{m}_{i \rightarrow j} = \mathbb{E}[\Delta^{-2} A_{0 \rightarrow \Delta}^2 \mid X_0 = i, X_\Delta = j]$ is the pair-conditioned second moment, so that $\bar{v}_{i \rightarrow j} = \bar{m}_{i \rightarrow j} - \bar{y}_{i \rightarrow j}^2$.

symbol	in the implementation	what it is
$P_{i \rightarrow j}(\Delta)$	P	transition over the window
$\bar{y}_{i \rightarrow j}$	gmean_ij	conductance averaged over the window, given the pair
$G_{i \rightarrow j} = \bar{y}_{i \rightarrow j} P_{i \rightarrow j}$	gtotal_ij	the same, weighted, so that it can be summed
$\bar{m}_{i \rightarrow j}$	gsqr_ij	its second moment, given the pair
$(\bar{y}_0)_i$	gmean_i	$= \text{gtotal_ij} \cdot \mathbf{1}$, Eq. 10
$\bar{m}_{i \rightarrow \cdot}$	gsqr_i	$= \text{gtotal_sqr_ij} \cdot \mathbf{1}$
$H_{i \rightarrow j}$, collapsed	$(\text{gtotal_ij} \circ \text{gmean_ij}) \cdot \mathbf{1}$	the square collapsed, not the collapse squared
$\bar{v}_{i \rightarrow j}$	gvar_ij	present in the window object, unused by the recursion
μ_0, Σ_0	P_mean, P_Cov	the state where the window opens
$\Sigma_0 - \text{diag } \mu_0$	SmD	the prior with its single-channel part removed
$\gamma^T \Sigma \gamma$ and its tilde	gSg	what the state contributes to the variance
$\mu^T \mathbf{w}$	ms	the state-dependent noise, Eq. 26
$\mathbf{g}_{0 \rightarrow \Delta}$	gS	state against current
$\sigma_{0 \rightarrow \Delta}^2$	y_var	the predictive variance

What is enough to reproduce this

The likelihood.

Every member of the family delivers, per window, a predicted mean $\bar{y}_t^{\text{pred}}$ and a predicted variance $\sigma_{\text{pred},t}^2$, and the recorded value is scored under the Gaussian they define,

$$\log q_t = -\frac{1}{2} \left[\log(2\pi\sigma_{\text{pred},t}^2) + \frac{\delta_t^2}{\sigma_{\text{pred},t}^2} \right], \quad \delta_t = \bar{y}_t^{\text{obs}} - \bar{y}_t^{\text{pred}}, \quad (36)$$

summed over windows. The recursion enters only through what $(\bar{y}_t^{\text{pred}}, \sigma_{\text{pred},t}^2)$ are conditioned on: a recursive member carries $(\mu^{\text{post}}, \Sigma^{\text{post}})$ of Eqs. 22 and 23 into the next window, so Eq. 36 sums to $\log \prod_t q_t(\bar{y}_t \mid \bar{y}_{1:t-1})$; a non-recursive member propagates from the initial condition by Eqs. 16 and 17 alone and the same sum is $\log \prod_t m_t(\bar{y}_t)$. Nothing else in this section is used by the fit.

The initial condition.

The recursion opens with the occupancy the model supplies and with the covariance that occupancy implies for one channel,

$$\mu_0 = \pi, \quad \Sigma_0 = \text{diag}(\pi) - \pi\pi^T. \quad (37)$$

For the schemes run here π is the first basis vector: the record opens at zero agonist with every channel shut, so $\Sigma_0 = \mathbf{0}$ and the spread the filter carries is generated entirely by the propagation of Eq. 17. Eq. 37 is also the one moment at which Σ genuinely is a single channel's covariance; from the first update onward, conditioning on a shared recorded value correlates the channels and what the recursion carries is $\text{Cov}(\mathbf{N})/N_{\text{ch}}$, which is no single channel's covariance. Every factor of N_{ch} in this section comes from that distinction.

What was verified, and to what tolerance.

The identities of this section were checked against a direct contraction over the K^2 pairs rather than asserted. Eqs. 20 and 21 reproduce the explicit contraction of the boundary covariance in 108 random cases to 7×10^{-18} and 3×10^{-16} ; the commutation of the start-index marginalization with the update, which is what licenses Eqs. 22 and 23 being written on the K states, holds in 36 cases to 2×10^{-16} ; Eq. 30 holds to 1×10^{-17} in 36 cases, independently of Σ_0 , and against a Monte Carlo of the continuous-time path; Eqs. 14 and 15 were checked against a Monte Carlo of 4×10^5 replicates of 200 channels, worst entry 2.9 standard errors over 256 entries; and the two forms of Eq. 26 agree with the forms that go through the pair-resolved variance to 1×10^{-15} in 16 cases. The whole construction, run on the pairs and marginalized, reproduces Eqs. 16, 17, 18, 19, 20, 21, 22 and 23 to 9.3×10^{-16} relative on two priors, a multinomial start and a perturbed post-update one.

The top of the ladder

The exact likelihood conditions on the full trajectory inside each window and is intractable, which is why the two closures above are needed at all. The same stochastic model that makes the exact likelihood intractable can be simulated forward exactly, one trajectory at a time, so the simulation realizes the object at the top of the ladder and serves as the ground truth against which every rung below it is judged.

Diagnostics: testing a likelihood against its simulation

An approximate likelihood is a misspecified model, and misspecification leaves a classical signature. At the true parameters the score no longer has zero mean, so the estimator the likelihood defines is biased; and the covariance of the score no longer equals the negative expected curvature, so the uncertainty the likelihood reports is wrong. These are the two Bartlett identities, and their failure is the content of the information-matrix-equality diagnostic of *White (1982)* and of the sandwich covariance of *Huber (1967)*; *Godambe (1960)*. In the classical setting those identities can only be tested, through a statistic with a known finite-sample defect, because the data-generating law is unknown. Our setting removes that limitation. The forward process is an exact continuous-time Markov chain (CTMC) and we simulate it exactly, so we hold unlimited data drawn from the very law the likelihood claims to describe, at parameters we set. The classical test therefore becomes a direct measurement: the anchor curvature is computed analytically and the score covariance is estimated by Monte Carlo over independent simulated recordings, with no null hypothesis to reject. The apparatus is standard; what is new here is applying it in this domain and reading its output as a measured property of each approximation rather than as a hypothesis to accept or reject.

We read three quantities off the simulation, in increasing strength, and each is read at a parameter vector that has to be named because there are two of them. The first is the simulation truth θ_{sim} , the point the recordings were generated from, which is known here and never known in an experiment. The second is the pooled optimum θ_{pool} , the maximum-likelihood fit of the member

under test to all the recordings of one cell at once; it is not known in closed form and is located numerically, by Gauss–Newton with Levenberg damping warm-started at the truth, with a near-zero gradient certifying that the point reached is stationary (Methods). Were the likelihood correctly specified the two would agree up to sampling error. They do not, and their difference $\theta_{\text{pool}} - \theta_{\text{sim}}$ is the misspecification bias itself. Everything below is computed at both anchors and every figure states which one it shows, because they answer different questions: at θ_{sim} a nonzero mean score is visible at all, whereas at θ_{pool} the score vanishes by construction, so the information comparison is not contaminated by a displaced gradient.

The first is the standardized residual. With μ_t and σ_t^2 the likelihood’s own predictive mean and variance for interval t , the residual $r_t = (y_t - \mu_t)/\sigma_t$ must have mean zero, unit variance, and no autocorrelation across intervals if the one-step predictive distribution is correct *at the parameters it is evaluated at* (Figure 3 draws the second and the third; the first is carried in a supplement, where it separates no member and so returns no verdict). This probes the predictive one interval at a time. It is the cheapest of the three and the only one an experimentalist can run on a real recording, because it needs no ensemble and no known truth; but it does need a fitted point, since at an arbitrary parameter vector the three properties fail for reasons that have nothing to do with whether the approximation is any good. Running it at the fit has a cost of its own. For a Gaussian likelihood the score is a weighted combination of exactly the quantities r_t and $r_t^2 - 1$, so setting it to zero imposes as many linear constraints on the residual moments as there are free parameters, and a fit that leaves the baseline and the noise level free absorbs most of the first two properties into itself. What no fit absorbs is the temporal structure. Residual whiteness is therefore the part of this check that stays informative at the optimum, and it is not optional: residual autocorrelation is the same signal that the strongest check quantifies below, though it is a conservative reading of it, running 1.6 to 1.7 times smaller than the score autocorrelation for the members in which either is large.

The second is the score. Writing $\ell(\theta)$ for the total log-likelihood and $S(\theta) = \nabla_{\theta} \ell(\theta)$ for its gradient, correct specification requires $\mathbb{E}[S(\theta_{\text{sim}})] = 0$ at the simulation truth. A nonzero mean score is what biases the estimate, and projected through the parameter covariance it becomes a displacement in parameter units, $b = \mathbf{H}^{-1} \mathbb{E}[S]$, with $\mathbf{H}$ the Gaussian Fisher information the likelihood itself reports (Eq. 38 below), the reader-facing statement of how far the estimate is pushed (Figure 3). This projection is a first-order approximation, from the expansion of the score at the optimum given below. Unlike the residual, this one cannot be run on an experiment at all, and for two reasons rather than one. It needs a known truth, and it needs an ensemble to take the expectation over. At the only point an experimenter has, their own fit, the score vanishes by construction and the check is empty. The simulator supplies both, and the confidence interval narrows as the number of simulated recordings grows.

The third and strongest is the information itself. Correct specification requires $\text{Cov}[S] = \mathbf{H}$, the information the likelihood reports, the covariance being taken over independent recordings of the same cell. That is the second thing an experiment cannot supply: one recording gives one score, and a covariance cannot be estimated from a single draw. An approximation can pass the residual and score checks and still violate this one by an order of magnitude. The two failures are independent: whether the estimate is biased is governed by how the conductance is modelled over the interval, while whether the reported information agrees with the information delivered is governed by the recursion. We turn all three checks on the whole cost ladder set out in Theory, from least squares on the deterministic mean current at the bottom, in both of its window settings, to the boundary-conditioned filter at the top. The range they have to cover is wide. The members that carry no recursion misreport the information grossly, and among the recursive members the interval conditioning leaves distortions small enough that a diagnostic has to be sharp to resolve them, so the same instrument has to remain informative across both. Because the failure is directional, we record it as a matrix,

$$\mathbf{C} = \mathbf{H}^{-1/2} \mathbf{J} \mathbf{H}^{-1/2},$$

with $\mathbf{J} = \text{Cov}(S)$ the covariance of the total score, cross-interval terms included. This definition is exact algebra given $\mathbf{H}$ and $\mathbf{J}$, and under correct specification $\mathbf{J} = \mathbf{H}$ so $\mathbf{C} = \mathbf{I}$ exactly. We fix the direction by convention: $C_{ii} > 1$ means the likelihood under-reports the uncertainty in parameter i and is over-confident, $C_{ii} < 1$ means it over-reports and is conservative. The eigen-directions of $\mathbf{C}$ say which combinations of parameters are misreported, which a scalar summary cannot recover.

The anchor $\mathbf{H}$ is the likelihood's own Gaussian Fisher information,

$$\mathbf{H} = \sum_i \mathbb{E} \left[\frac{(\nabla \mu_i)(\nabla \mu_i)^\top}{\sigma_i^2} + \frac{(\nabla \sigma_i^2)(\nabla \sigma_i^2)^\top}{2\sigma_i^4} \right], \quad (38)$$

formed from the same predictive moments and their parameter derivatives that the score already requires, at no extra cost. For a macroscopic algorithm the likelihood is Gaussian by construction, so this is the model's own Fisher and $\mathbf{C}$ is the information-matrix-equality test carried out in the model's own frame. Two properties follow. The Gaussian Fisher is positive semidefinite by construction, so the diagnostic never has to be rescued from a numerically indefinite anchor. And it is analytic, so it does not inherit the finite-difference noise of a numerical Hessian, which in this regime is severe: the finite-difference Fisher is widely indefinite per replicate, and only pooling across replicates recovers a definite matrix. By design the anchor is the Gaussian Fisher; the numerical Fisher we compute alongside it serves to gauge how faithful that Gaussian anchor is. Comparing the approximation to its own Fisher is not circular. $\mathbf{H}$ is the information the likelihood claims, $\mathbf{J}$ is the information it delivers against data from the true process, the two are computed from different objects, and their disagreement is the definition of misspecification.

The bottom rung of the ladder needs a qualification at the point where the diagnostic is defined, because its predictive variance is not a gating quantity. Least squares fits the deterministic mean current under a single variance held constant over the whole recording, so the information it reports, $\mathbf{H}$ as written above with σ_i^2 replaced by one constant σ^2 , is the information of a homoscedastic model. This holds of both least-squares arms: averaging the predicted mean over the acquisition interval changes μ_i and leaves the constant-variance premise untouched. The exact process violates that assumption: the variance of a recorded interval rises and falls with the open population over the course of the concentration jump, and it also depends on the channel number, so the constant-variance premise is wrong by construction in every cell we simulate. For a least-squares arm the quantity $\text{Cov}(S)$ read against the reported information therefore measures the size of that break. It is not an information-matrix-equality failure of the same kind as a gating-aware likelihood's, where the predictive variance does follow the gating and the discrepancy is attributable to how the gating is approximated within the interval and across intervals. The Results quote the least-squares ratio with that reading, and they do not read it and a gating-aware ratio as two values of one comparable quantity.

The same matrix repairs what it measures. The sandwich covariance

$$\Sigma = \mathbf{H}^{-1} \mathbf{J} \mathbf{H}^{-1} = \mathbf{H}^{-1/2} \mathbf{C} \mathbf{H}^{-1/2}$$

is the corrected parameter covariance: it substitutes the information the likelihood actually carries for the information it reported. This step, and the bias vector b above, are first-order approximations, from a Taylor expansion of the score at the optimum with the empirical Hessian replaced by its expectation $-\mathbf{H}$ (see Supplementary Information). We verify the correction directly, without leaning on the expansion, by comparing the sandwich-corrected ellipse against the empirical covariance of the maximum-likelihood estimates (MLEs) over the simulated recordings (Figure 2). A diagnostic that also corrects the error bars it flags is a stronger result than one that only scores algorithms.

The distortion splits into two physically distinct parts (Figure 4-figure supplement 1). The sample distortion $\mathbf{C}_s = \mathbf{H}^{-1/2} \mathbf{J}_s \mathbf{H}^{-1/2}$, built from the within-interval score covariance $\mathbf{J}_s = \sum_i \text{Cov}(s_i)$ with $s_i = \nabla_{\theta} \ell_i$ the score of interval i , measures per-sample departure from Gaussianity; to the leading corrections it is the third and fourth cumulants of the interval current, the non-Gaussian shape

the Gaussian predictive cannot carry. The correlation distortion $\mathbf{R} = \mathbf{J}_s^{-1/2} \mathbf{J} \mathbf{J}_s^{-1/2}$ measures what the within-interval covariance misses, the cross-interval correlation of the score, which is exactly the local time correlation of the current that *Milescu et al. (2005)* named as the dominant error of non-recursive fitting; $\mathbf{R} = \mathbf{I}$ when the score carries no cross-interval dependence. The figures resolve the same pair in time and one parameter at a time, and the notation follows the matrices. For the parameter under discussion, F_t is the Gaussian Fisher information of interval t alone, so that the corresponding diagonal entry of $\mathbf{H}$ is $\sum_t F_t$; $J_t = \text{Var}(s_t)$ is the variance of that interval's score. Accumulated to interval T they are $F_T = \sum_{t \leq T} F_t$ and $J_T = \text{Var}(\sum_{t \leq T} s_t)$, and it is the second of these, not the sum of the J_t , that the reported information has to match. The two statements $J_t = F_t$ and $J_T = F_T$ are therefore not equivalent, and their divergence is the whole of what Figure 3 measures. These compose exactly. The exact identity is

$$\mathbf{C} = \mathbf{K} \mathbf{R} \mathbf{K}^\top, \quad \mathbf{K} = \mathbf{H}^{-1/2} \mathbf{J}_s^{1/2},$$

which holds as algebra for any $\mathbf{H}$ and $\mathbf{J}_s$, since substituting the definitions collapses the chain. The factor $\mathbf{K}$ satisfies $\mathbf{K} \mathbf{K}^\top = \mathbf{C}_s$, so it is a legitimate square-root factor of the sample distortion, but it is not the symmetric principal square root, and the form $\mathbf{C} = \mathbf{C}_s^{1/2} \mathbf{R} \mathbf{C}_s^{1/2}$ holds only where $\mathbf{H}$ and $\mathbf{J}_s$ commute. What composes with no assumption at all is the information volume,

$$\log \det \mathbf{C} = \log \det \mathbf{C}_s + \log \det \mathbf{R},$$

an exact identity because the determinant is multiplicative. The per-parameter diagonal, the trace, the eigenvalues and the Frobenius norm do not compose across the split; only the volume does, and only the volume is decomposed in the maps.

The correlation distortion has a reader-facing translation. For any additive statistic X_t with sum $L = \sum_t X_t$, the ratio $\kappa = \text{Var}(L) / \sum_t \text{Var}(X_t)$ is unity when the terms are independent and larger when they are positively correlated, and it converts to an effective sample size $T_{\text{eff}} = T/\kappa$: a non-recursive likelihood believes it has T independent intervals when it effectively has T/κ . This is the number from the diagnostic most worth quoting, and it is exact by definition. When $\mathbf{R} \approx \kappa \mathbf{I}$ the temporal error is isotropic and this scalar captures it; when $\mathbf{R}$ is anisotropic, or when $\mathbf{C}_s \neq \mathbf{I}$, the distortion cannot be reduced to a sample-size correction and the full matrix is needed.

A matrix is not a number, and two scalars summarise it without pretending to replace it. Let $\lambda_1, \dots, \lambda_d$ be the eigenvalues of $\mathbf{C}$ on the retained subspace. Their logarithms have a mean $\bar{e} = d^{-1} \sum_i \log \lambda_i$ and a variance $s^2 = \text{Var}_i(\log \lambda_i)$, and the two quantities reported throughout are their exponentials,

$$m = e^{\bar{e}} = \left(\prod_i \lambda_i \right)^{1/d}, \quad a = e^s, \quad (39)$$

so that m is the typical factor by which the reported information is wrong, the geometric mean of the per-direction factors, and a is the multiplicative spread of those factors from direction to direction. A member is drawn as a point at m with a bar from m/a to ma , which is why the axes are logarithmic and why the spread is written as a factor rather than as a plus-or-minus. Both are one at correct specification. The single scalar that combines them is the affine-invariant Riemannian distance from the identity, $D_{\text{AI}} = \left(\sum_i \log^2 \lambda_i \right)^{1/2} = \sqrt{d(\bar{e}^2 + s^2)}$, which is used where a single number is wanted and is reported nowhere in place of the pair. In error-bar units take the square root of any of these, since the distortion is a ratio of informations: a distortion of 1.15 is a 7% error on a confidence interval, 2 is 41%, and 4 turns a nominal 95% interval into about 68%.

Two conventions govern the edge cases. The direction is fixed as above, $C_{ii} > 1$ being over-confidence and $C_{ii} < 1$ conservatism, stated once so every verdict reads the same way. And the anchor can be near-singular where a parameter is unidentifiable, which by design happens once the open population relaxes (Figure 3-figure supplements 1 and 2); there we diagonalize $\mathbf{H}$ and retain the eigenmodes whose eigenvalue exceeds a tolerance set relative to the largest, $\lambda_i > \lambda_{\text{max}} \cdot 10^{-10}$, evaluate every inverse power on that retained subspace alone, and leave the quantity undefined, not finite, when the retained subspace is empty. The dimension d of Eq. 39 is the dimension of

that subspace, so a cell where a direction has been dropped is summarised over fewer directions than one where none has. Undefined is not zero, and the maps render those cells in grey to keep the two apart.

Results

A recursive likelihood cannot be checked against a record by eye. Least squares predicts each sample from the parameters alone, so the prediction can be laid over the data and judged where it misses. A filter predicts each sample from the samples before it, so it follows the record whether or not the uncertainty it reports is the uncertainty it delivers. What follows measures the second against the first, on recordings simulated exactly from the process every member approximates.

Throughout, N_{ch} is the number of channels in the patch, k_{on} and k_{off} the opening and closing rates, i the unitary current, Δ the acquisition interval and $\tau = 1/k_{\text{off}}$ the closing time constant. A tilde marks a quantity in the natural units of the model, time in τ and current in i : the acquisition interval is reported as $\tilde{\Delta} = \Delta k_{\text{off}} = \Delta/\tau$, and the instrumental noise as $\tilde{S}$, the instrumental noise power accumulated over one channel time constant measured against the single-channel signal power (Methods). The two least-squares arms are one flag apart, LSE fitting the conductance at the sample and ILSE fitting it averaged over the acquisition window, and both are measured. Their coverage differs: ILSE is run over the whole design plane, and LSE has the clouds of the single-cell comparison at three channel counts and the distortion measurements across the plane.

Recovery at one design cell

Figure 2 turns the ladder into a statement about reported uncertainty at one cell: $N_{\text{ch}} = 100$, $\tilde{S} = 0.01$, $\tilde{\Delta} = 0.1$. Ten thousand independent recordings are partitioned into groups of ten, each group is fitted by maximum likelihood under each member, and the thousand estimates that result are compared against the covariance that member reports from its own Gaussian Fisher information. Every panel carries three readings: where the cloud sits, how large the reported ellipse is, and what shape it has.

The first reading separates the family by the interval average, and the roster isolates that step twice, once in each family and one flag apart both times: LSE against ILSE, and NR against INR. In the kinetic pair every member lands within 0.04 in $\log_{10}$ of the truth, and both times the interval average takes a factor of two to three off the offset in k_{on} , from -0.015 to -0.007 in the least-squares family and from -0.021 to -0.008 in the macroscopic one. Recursion on its own does the opposite: the recursive instantaneous filter is the least centred member in the pair, at -0.040 , worse than the open-loop member it is built from, and the one-endpoint rungs carry -0.031 and -0.012 . Only IR, which has both, has no offset resolved from zero in either coordinate.

In the amplitude directions the same split costs an order of magnitude more, and there the displacement has a mechanism the panel can be read for. The unitary current comes out low by 0.228 for MR, 0.156 for R, 0.117 for VR and 0.103 for NR, which is a factor of 0.59 at the worst, against 0.028 for INR and nothing resolved for IR; the channel number comes out high by almost exactly the same amount in each case. What the mean current determines is only the product $N_{\text{ch}} i$, and every member recovers that product to within 0.023 in $\log_{10}$. What separates the two factors is the gating variance, whose magnitude carries a second power of the unitary current, so a member with the wrong interval variance splits a well-determined product wrongly and misses both factors by a large amount with opposite signs.

The noise level runs the same way and states it most sharply. NR carries no interval-mean conductance variance, so the gating variance it predicts is too small, and the fit makes up the difference in the one free parameter that can absorb it: its noise estimate lands 0.333 in $\log_{10}$ above the truth, a factor of 2.2, against every other member within 0.04. Conditioning the conductance on the interval instead of on an instant, and adding nothing else, returns it to 0.005. Modelling the acquisition average therefore buys the noise level, and it buys it with no recursion at all.

The open marker in each panel is the truth displaced by the bias the distortion theory predicts, and that comparison is partial. On k_{off} the prediction falls inside the sampling interval of the measured displacement for all eight members. In the two amplitude directions it carries the sign and between 71% and 85% of it for NR, R, MR and VR, and it over-states NR's noise offset by a factor of 1.8. The first moment is diagnosed to within a fifth rather than reproduced, which is what a first-order prediction evaluated at the truth can be expected to do for a cloud fitted in groups of ten, and the quantitative claim of this section is the second moment.

The second reading is size. The reported covariance disagrees with the cloud by an order of magnitude at the foot of the ladder and by nothing at its top. Taking the ratio of the area of the empirical 95% ellipse to the area of the reported one, and averaging geometrically over the six pairs the four informative parameters form, the open-loop member reads 11.41, the recursive instantaneous member 1.17 and the boundary-conditioned member 1.01; on the joint four-parameter ellipse the same three read 154.65, 1.56 and 1.12. A likelihood that carries no recursion under-reports the area of the parameter distribution by a factor of ten to fifteen, and the recursive members are wrong by tens of percent.

The fall between them is not monotone, and this is where the ladder stops being an ordering. On the kinetic pair the recursive instantaneous filter reads 1.32 while MR reads 1.97 and VR reads 2.19, and only then does IR return to 1.02; the amplitude and noise pairs run the same way. Conditioning the conductance on one endpoint costs calibration rather than buying it, in every pair and in both directions of the variance switch, and what recovers it is the second endpoint entering the gain (Appendix 1). The two intermediate rungs are kept in the figure for that reason: they are the step from R to IR taken in its two algebraic halves, and the half that changes the interval variance moves the answer the wrong way.

What the acquisition average buys on its own is separable here, because one member isolates it. The one-endpoint non-recursive member conditions the conductance on the interval instead of on an instant and adds nothing else, and against the open-loop member it improves every pair, most on the noise pair, from 8.84 to 3.28. The same step taken inside the least-squares family buys nothing at all in the one pair both arms carry: the kinetic area ratio is 14.73 for LSE and 14.77 for ILSE, while their first-moment offsets differ by a factor of two. Averaging the model over the acquisition window is worth a great deal to a likelihood that also models the gating variance and next to nothing to one that does not, and that last improvement, the one from 8.84 to 3.28, is the one the third reading takes apart.

The third reading is shape, and it is why a size is not a diagnosis. The thin dot-dash ellipse in every panel is the reported one resized to the empirical area with nothing else touched, which is the most any single factor applied to an error bar can do, and it does not land on the empirical ellipse. The second orange number grades what is left over, $\sqrt{\lambda_{\text{max}}/\lambda_{\text{min}}}$ for the empirical covariance against the reported one, so that the two orange numbers are the two eigenvalues of the distortion written as a geometric mean and a spread and the worst and best directions are their product and their quotient. It runs from 1.0 to 3.2 across the panels against a finite-sample floor of at least 1.08, that floor being the 95th percentile of the same statistic when a thousand estimates in two dimensions are drawn from a perfectly calibrated covariance, and a lower bound because the reported covariance is itself estimated.

The clearest case is INR in the noise pair, where the two numbers read 3.3 and 3.2. They factor into a variance ratio of 10.5 along one direction and 1.03 along the other, and the calibrated direction is the noise level itself, within a tenth of a degree of that axis. The member that repairs the noise estimate reports the noise uncertainty exactly right and the channel-number uncertainty with a standard error 3.2 times too small, and the single factor a scalar correction would apply, 3.3, is right for neither of them. The alignment is not particular to that member: in the same pair MR and VR read 1.04 and 1.02 along the noise axis against 1.59 and 2.02 across it. One panel is an instance; the same anisotropy holds over the design plane, and the repair it forecloses is taken up with it in Figure 4-figure supplement 3.

The distortion is measurable and it is also removable, provided the correction is a matrix. Replacing each reported covariance by the sandwich returns every gating-aware member of the ladder to within about a tenth of one, with every parameter pair between 0.87 and 1.07, and it is the least-squares family alone that it leaves short, at 1.65 and 1.63 for the two arms. The two arms fail it by the same amount while differing in every first-moment quantity, which puts the failure in what they share and not in the interval average. The reader-facing form of that repair is coverage. Reducing the estimates to their squared Mahalanobis distance from the cloud mean, which is χ^2 with as many degrees of freedom as free parameters if the reported covariance is right, the nominal 95% region of IR covers 0.914 of them under its own Fisher covariance and 0.947 to 0.949 under the sandwich, at $N_{\text{ch}} = 10$ and $\tilde{S} = 0.005$ over a thousand fits in six parameters at group size 100 (Figure 4–figure supplement 2). That one correction repairs six of the eight members, without being told which of them is wrong or how, is what makes the distortion a usable quantity rather than a complaint.

Both least-squares arms carry a column in the kinetic pair and none in the other two, for reasons of identifiability rather than coverage: their grid configuration holds the unitary current and the noise level fixed at the simulated values (Methods), so in each of those pairs one axis has no spread, both ellipses collapse to a segment and the area ratio would be 0/0. Each does estimate the channel number, and to within 0.012 in $\log_{10}$, but with the unitary current pinned at its true value that estimate is the mean current under another name. What this figure measures the members on is the separation of the two factors, which is the harder half and the one the gating variance carries.

The calibration cascade in time

The miscalibration of Figure 2 is a property of the likelihood and not of the optimiser, and Figure 3 shows where in the recording it is produced. One common set of a thousand recordings at the same cell is scored interval by interval by each member, with no fitting involved. The least-squares arm here is the six-parameter configuration, which keeps the unitary current and the noise level free so that its per-interval quantities align index for index with the macroscopic dumps (Methods); nothing measured in this subsection may be carried across to the four-parameter configuration of the later figures.

The members separate first on the log-likelihood they assign to the same data, over a span of 118.9 nats from the bottom of the ladder to the top, and most of that span is bought by modelling the gating fluctuations at all.

One pair on the ladder does not separate at all. The two least-squares arms assign the same log-likelihood to the same data, -262 nats, and agree to three significant figures on every statistic the figure reports, the accumulated information ratio at 14.96 and the lag-one score autocorrelation at 0.856 among them. Carrying one constant variance and no occupancy distribution, the classical fit gives the acquisition average nothing to act on except the predicted mean itself, and at this cell that displacement is too small to register on any of the diagnostics. The same step taken inside the family moves every one of those numbers, from NR to INR and from R to IR. The window axis is defined at all three levels of the gating description and is only consequential at two of them.

The hinge of the argument is that per interval every member reports its information correctly. The variance of the per-interval score matches the per-interval Fisher information, with the median of $\log_{10}(J_i/F_i)$ over the informative steps inside ± 0.19 of zero for ILSE, NR, R and IR on every parameter. Accumulated over the recording the same ratio reads 14.96 for either least-squares arm and 13.63 for NR, against 1.07 and 1.08 for the two recursive members. The open-loop members are within a few percent of the identity at each interval and off by more than an order of magnitude once the intervals are summed. The open-loop member also changes the sign of the discrepancy between the per-step and the accumulated statistic, and a per-sample modelling error cannot do that.

The least-squares ratio needs the qualification the Diagnostics section attaches to it. The Gaus-




Figure 2. Parameter recovery and the calibration of the reported interval, for the whole ladder at one design cell. $N_{\text{ch}} = 100$, $\tilde{S} = 0.01$, $\tilde{\Delta} = 0.1$; 10,000 recordings in groups of ten, so one thousand maximum-likelihood estimates per member. Columns are the eight members, in ladder order, from least squares to the boundary-conditioned filter. The first two are the least-squares arms, LSE fitting the conductance at the sample and ILSE fitting it averaged over the acquisition window, and the pair NR, INR is the same one step taken inside the macroscopic family. Rows are three parameter pairs, each chosen because a different property separates the members in it: the kinetic pair ($\log_{10} k_{\text{on}}$, $\log_{10} k_{\text{off}}$), which the deterministic transient determines and on which every member including both least-squares arms competes; the amplitude pair ($\log_{10} N_{\text{ch}}$, $\log_{10} i$), which only the gating fluctuations separate; and the noise pair ($\log_{10} \sigma$, $\log_{10} N_{\text{ch}}$), where σ is the instrumental noise level, a free parameter of every likelihood member, with a simulated value of 10^{-4} . Each panel carries the cloud of estimates and four 95% ellipses drawn at the cloud mean: the empirical covariance, the Fisher covariance the member reports, that same Fisher ellipse resized to the empirical area and otherwise untouched (thin, dot-dash), and the sandwich covariance. The blue number is the area of the empirical ellipse over the area of the corrected one. The two orange numbers are the same area ratio against the reported ellipse and, after it, the spread $\sqrt{\lambda_{\text{max}}/\lambda_{\text{min}}}$ of the reported covariance against the empirical one, written $\times$: they are the two eigenvalues of the distortion as a geometric mean and a spread, so a member is calibrated in size when the first reads 1.0, calibrated in shape when the second does, and the resized ellipse falls on the empirical one only in the second case. Markers are the simulation truth, the mean of the cloud, and the truth displaced by the predicted bias. The four least-squares cells that are empty are empty by construction and say so: that configuration holds the unitary current and the noise level fixed at their simulated values, so those estimates would lie on a line and the area ratio would be 0/0. Covariances are the Gaussian-Fisher family at the pooled estimate, scaled by the reciprocal of the group size. What the third row separates is in the text.

sian Fisher information of a least-squares fit presumes one constant residual variance, and the exact simulator violates that premise by construction, because the variance of a recorded interval rises and falls with the open population. The factor of fifteen is the size of that break, and it is not an information-matrix-equality failure of the same kind as the open-loop gating likelihood's, where the predictive variance does follow the gating. The two ratios are reported side by side and are not read as two values of one comparable quantity.

The mechanism is visible in the same figure. At lag one the autocorrelation of the per-interval score on k_{off} is -0.0043 ± 0.0043 for IR, indistinguishable from zero, against 0.191 ± 0.004 for R, 0.864 ± 0.0019 for NR and 0.856 ± 0.0022 for either least-squares arm. A positively correlated score makes the variance of its sum exceed the sum of the per-interval informations, which is how the per-interval identity holds while the accumulated one fails. The temporal correlation the two non-recursive members leave in their score is the local time correlation of the current that *Milescu et al. (2005)* named as the dominant error of independent-interval fitting, measured here at the level of the inference rather than the data.

Row C of Figure 3 carries a second result. For the members that condition on the past, the information about the channel number and the opening rate reaches the numerical floor within six intervals of the agonist being removed: having conditioned on everything up to that point, the filter gains nothing further from watching the channels decay. Two parameters survive the washout by different routes. The closing rate survives through the mean, which is the shape of the decay its own kinetics govern, and the unitary current through the variance, which is the gating fluctuation of the channels still open. This is a property of conditioning rather than of the macroscopic observable: the two non-recursive members retain 15 to 16% of their pulse information about the channel number through the washout, against 0.5 to 0.8% for the two recursive ones. The parameters the figure has no colours for, and the first moment of the residual, are in Figure 3-figure supplement 3; the instrumental noise level is the one parameter measured before the agonist arrives, and its score is white even in the members whose kinetic directions are not, so the memory that carries the accumulated error is a property of the gating directions.

The design plane, in both moments

One cell is not a result, and the plane carries one the cell cannot: instrumental noise equalises the members. Where it grows past the gating variance every rung reports honestly, so the differences between them are a property of the regime in which the fluctuations still carry information, and not of the algorithms in general. Everything below is read against that.

Figure 4 repeats the two moments over the whole grid, with the members across the columns, the channel count down the rows, and the acquisition interval against the instrumental noise inside each panel. It splits the display by moment rather than by member, so the least-squares arm stays alongside the likelihoods it is being compared with. Half (A) is evaluated at the simulation truth, where a bias is visible; half (B) at the pooled optimum, where the score vanishes by construction and the information comparison is clean (Methods).

More channels do not repair the recursive instantaneous filter. Taking the median over the seven intervals at $\tilde{S} = 0.01$, its k_{off} distortion sits on a flat floor near 1.4 across four decades of channel count while the boundary-conditioned member converges to one over the same range; the open-loop member moves the other way, from 22 to 77, and the least-squares arm holds a floor near 13. In the channel-number direction the two recursive members separate in sign as well as in size: the boundary-conditioned one stays on the conservative side throughout, while the recursive instantaneous one crosses from conservative to over-confident and keeps rising.

The first moment carries a separate cost. The recursive instantaneous filter has a bias in the channel number of about 0.14 in $\log_{10}$ units, close to a 35 to 40% error that does not shrink as channels are added; the open-loop member carries half of that in the same direction, and the boundary-conditioned member is at or near zero everywhere. Least squares has a rate bias of essentially zero over the same plane, and that concession is worth making plainly: the classical fit




Figure 3. The calibration cascade in time, over the whole implemented family. One set of 1,000 recordings ($N_{ch} = 100$, $\hat{S} = 0.01$, $\hat{\Delta} = 0.1$) is scored by each of the eight members (columns, in cost-ladder order), with no fitting. The two least-squares arms open the row, LSE on the instantaneous mean and ILSE on the interval-averaged one. Rows A and B are properties of the data, in grey; rows C to G are per parameter, the closing rate and the channel number in colour. (A) An example recording with the predicted current, its band the filter's own predictive ± 1 s.d., annotated with the ensemble-mean log-likelihood and its standard error. (B) The mean squared standardized residual, one when the predictive variance is right. (C) The per-interval Fisher information, so where each parameter is measured. (D) The mean per-interval score over $\sqrt{F_t}$, zero when unbiased and standardized so parameters with different units share an axis. (E) The per-interval information ratio $\text{Var}(s_t)/F_t$, one at each interval. (F) The same ratio accumulated, J_T/F_T . (G) The autocorrelation against lag of the standardized residual, in black, and of the score for each parameter, in its colour; the score's is why (E) holds where (F) does not, and the residual's is the same memory in the data and the only quantity here an experimentalist can compute on a real recording. Bands are 95% bootstrap intervals over the recordings; curves stop where the information reaches zero. Non-negative quantities are on a logarithmic axis with the calibrated value on the ticks, signed quantities on a linear axis with zero on the ticks; the four open-loop columns share a y-scale block distinct from the four recursive ones. The channel number is not drawn for either least-squares arm, since neither separates it from the unitary current: the mean current depends on the two only through their product, and it is the gating variance, absent from both, that breaks the tie.

Figure 3—figure supplement 1. Per-step information against score variance, for the opening rate and the closing rate.

Figure 3—figure supplement 2. The same, for the unitary current and the channel number.

Figure 3—figure supplement 3. The checks this figure has no room for: the mean standardized residual, the score bias in the three parameters not coloured above, and their score autocorrelation.

recovers the rates, and what it gets wrong is how well it says it recovered them.

Read together, the two moments say that the two axes of the family act on different failures and that neither substitutes for the other. Averaging over the acquisition interval is what removes the amplitude bias: the median absolute bias in the channel number is 0.099 and 0.110 in $\log_{10}$ units for the two members that do not average, against 0.003 and 0.002 for their interval-averaged counterparts, so recursion alone does not remove it and slightly worsens it. Recursion is what removes the distortion: the same step drops the closing-rate distortion from tens to about 1.4, and it reaches one when the recursion conditions on both ends of the interval. The same factorial appears again at one cell and resolved in time, in the four corner values of Figure 3: interval averaging alone leaves the per-interval residual calibrated and the accumulated information ratio at 20.8, recursion alone inverts that pair to 0.782 and 1.07, and the member with both is at one on each. A member that takes one axis and not the other repairs one failure and keeps the other, which is why the family cannot be ordered on a single scale. Nothing in this statement is specific to ion channels: it holds for any state model whose observations are averages over a window.

The two recursive members also differ in where on the plane their departures sit. Counting the fraction of cells that depart from one by at least 15%, parameter by parameter, the recursive instantaneous member departs on between a fifth and a half of the plane in four of the five directions, against at most a twelfth for the boundary-conditioned one. The distortion of the instantaneous filter is spread across parameters and across the plane; the distortion of the boundary-conditioned filter is confined to the few-channel corner. Conditioning localises the error as well as shrinking it.

Instrumental noise moves the whole picture continuously, and it moves the window-ignoring members furthest. The closing-rate distortion of least squares falls from its floor near 13 to one over about three decades of added noise, at every channel count, and the recursive instantaneous filter follows the same pattern about three decades lower. The reason is that instrumental noise adds a variance that carries no gating memory, so as it grows it dilutes the temporal structure a window-ignoring likelihood fails to model, until there is nothing left to fail on. The channel count at which each member's error bar becomes honest therefore rises with the noise level, which is the crossover Figure 6 draws as a boundary.

Over the whole measured grid, counting every parameter at every cell, the reported uncertainty is within $\pm 15\%$ of the delivered one for 93% of the boundary-conditioned member's points, 70% of the recursive instantaneous member's and 31% of the open-loop member's. Least squares carries no entry in this comparison, because its distortion matrix has four dimensions against the likelihood members' six and the two counts are not commensurable (Methods). The coverage is unequal and has to be quoted with the number: the boundary-conditioned member spans eleven channel counts from 5 to 10^4 across the twelve noise levels of the ladder (Methods), not on a full rectangle, and the others cover less.

The mechanism behind the flat floor is the split of the distortion into a per-sample part and a correlation part (Figure 4–figure supplement 1). On the closing rate at $\tilde{S} = 0.01$, the per-sample part goes to one as channels are added for both recursive members, 1.09 to 0.998 for R and 1.10 to 1.007 for IR across the four decades, so the single-interval Gaussian approximation becomes exact for both at the same rate. What separates them is the correlation part, which rises for R from 1.35 to 1.47 and falls for IR from 1.24 to 1.00. The two members are equally faithful to a single interval, and they part company on how much temporal dependence they leave in the score, which is what conditioning on both interval endpoints removes.

Two features of the plane are easy to misread. Where the acquisition interval is far shorter than the channel's correlation time, successive samples carry redundant information, and a faithful likelihood has to register that redundancy, so a large correlation distortion at the shortest intervals is the correct cost of oversampling rather than a defect. And at $N_{\text{ch}} = 10^4$ with $\tilde{\Delta} = 1$ the amplitude directions of the recursive instantaneous filter are not identified: the Gaussian Fisher information is ill-conditioned there, so a displacement that costs almost no likelihood is amplified into a large parameter shift. Those cells are drawn grey. Their sign and their existence are real; their magnitude

is a property of the conditioning.

Whether one correction can undo it

A reader who accepts the distortion will reach for the cheapest repair before changing likelihood. If a likelihood over-states its information, declare that the recording holds fewer independent intervals than it appears to and deflate the count. The number that repair would use is already measured: the integrated autocorrelation of the standardized residual, which counts how many intervals a likelihood effectively has for every interval it believes it has, reads 1.008 for IR, 1.30 for R, 3.93 for NR and 9.26 for ILSE at $\tilde{S} = 0.01$. The classical fit therefore behaves as though it had about nine times more independent observations than the recording supplies, which is the degrees-of-freedom cost of fitting correlated data as though it were independent, and it is a quantity an experimentalist can compute on a real recording without an ensemble and without a known truth.

The repair works if the distortion is a scalar, every parameter direction wrong by the same factor. Figure 4–figure supplement 3 measures whether it is, by splitting the distortion matrix into the two invariants of its spectrum: the geometric mean of the eigenvalues, which is the typical factor by which the reported information is wrong over all directions at once, and the spread of those eigenvalues about it, which is how much that factor varies from direction to direction. The two are the orthogonal components of the affine-invariant distance the diagnostic already reports, so a spread of one means a single rescaling is exact and the factor it should use is the magnitude.

The spread is not one where it would matter. The least-squares arm and the open-loop interval member each over-state their information by a factor of about 5 to 6, and the per-direction factor spreads by about 3.2 around it, flat across four decades of channel count, so directions exist that are wrong by twenty and directions exist that are barely wrong at all. No single effective count serves both. The claim a reader would otherwise make, that the residual autocorrelation supplies the correction, is refuted by the same measurement that would have supplied the number, and the two boundaries of the map below are therefore boundaries on which likelihood to use and not on which correction to apply.

The three recursive members separate on different axes of the same pair. The recursive instantaneous member's magnitude drifts up with the channel count while its spread sits flat at about 1.18, so its failure is mostly a size it could in principle correct for. The boundary-conditioned member is at one in magnitude at every count and its spread falls toward one, so it is the member converging on both axes at once; at ten channels its spread is 1.25, above the strict criterion, which is the same few-channel corner every other measurement in this section points at. The one-endpoint member is worse than the member it improves on, on both axes and at every channel count, which is the ladder claim in its sharpest form: a partial interval correction does not interpolate between the member without one and the member with a complete one.

Where the calibrated member departs, and which moment carries it

Figure 5 resolves the corner where the boundary-conditioned member departs, on a scale that Figure 4 cannot show: there a factor of 1.5 is a faint tint, and here it needs an axis. Finding that corner was half the object of the measurement, since a diagnostic that never fires on the member it is applied to says nothing about what it can resolve.

The departure is two-sided within a single cell. At ten channels and $\tilde{S} = 0.005$ the reported error bar of the closing rate is too small by a factor of 1.62 at the shortest interval, while that of the channel number is too large by a factor of 0.65, so one likelihood is over-confident about a rate and conservative about a count at the same time and in the same recording. The two parts of the decomposition answer to the acquisition interval in opposite directions. Across that same cell the correlation part of the closing rate falls from 1.52 to 1.27 as the interval coarsens, while the per-sample part rises from 1.06 to 1.19: sampling more coarsely buys back the memory and pays




Figure 4. The design plane: how wrong each member's report is, in both moments, on one scale. Ten thousand recordings are simulated at every cell of the grid and scored by each member of the cost ladder. (A) The distortion-induced bias, evaluated at the simulation truth θ_{sim} : how far the estimate is displaced. (B) The information distortion, evaluated at the pooled optimum θ_{pool} , where the score vanishes by construction so a displaced gradient cannot contaminate it: how far the uncertainty a member reports is from the one it delivers. Columns nest twice, member and then parameter, k_{off} and N_{ch} . The least-squares member on the grid is ILSE, the interval-averaged arm, and it holds one column per half, because its four-parameter configuration fixes the unitary current and the noise level at the simulated values (Methods), so a channel-number comparison against it would not be like for like. Rows are the channel count, each given a height proportional to the range of instrumental noise the sweep reaches there, so a decade of noise is the same distance in every row. Inside a panel the axes are the sampling interval in units of τ and the dimensionless instrumental noise, and one decade is the same physical distance on both, so a field's shape can be read as a shape. Colour is shared by the two halves: the factor by which the report is wrong, centred on one, pale right and saturated wrong in either direction, blue conservative. What a factor means differs between them and the difference is not cosmetic: in (A) it is a factor on the parameter, so 2 means the estimate is out by a factor of two, while in (B) it is a ratio of variances, so 2 means the reported error bar is $\sqrt{2}$, about 40%, too narrow. The scale is clipped at $10^{\pm 3}$, which touches 22 of 3276 cells, twenty of them the open-loop member, where being wrong by 10^3 and by 10^{12} say the same thing. Lines carry two codes: colour says which quantity, dark hairlines cutting the colour field and white lines on a dark casing cutting the distortion-corrected standard error, which asks whether the parameter can be measured at all rather than whether its error bar is honest; dash says which criterion, solid a factor 1.15 and dashed a factor 2, the two thresholds used throughout. Both criteria are edges of the colour bands, which is why the dark pair needs no key of its own.

Figure 4—figure supplement 1. The distortion decomposed into a per-sample part and a correlation part, against channel count and resolved by instrumental noise.

Figure 4—figure supplement 2. The sandwich prediction against the empirical multivariate distribution of the estimates, as a Mahalanobis quantile-quantile plot.

Figure 4—figure supplement 3. Can the distortion be undone by counting fewer intervals? The size of the information distortion against its shape.

for it in the per-sample non-Gaussianity. Both parts vanish by $\tilde{S} = 1$, where the instrumental noise has made the emission Gaussian enough that neither closure is strained.

Two limits travel with the figure. The family applies both of its Gaussian approximations at every step and the score sees only their product, so a departure measured here cannot be assigned to one of them; separating them needs the microscopic recursion, which is a companion paper's subject. And the grid is truncated where the effect grows: the departure of the closing rate is still accelerating at the lowest simulated noise, reading 1.22, 1.37 and 1.52 over the last three noise steps at ten channels, so the figure locates the onset and not the far side.

A usage map for a macroscopic recording




Figure 5. Where the boundary-conditioned filter stops reporting an honest error bar, and which of the two discarded moments is responsible. The same measurement as Figure 4, the Gaussian information distortion at the pooled optimum on the same runs, resolved on the corner where IR departs: on Figure 4's scale a factor of 1.5 is a faint tint and here it needs an axis. The three blocks are the distortion and the two parts it decomposes into: (A) the total, (B) the part carried by the per-sample non-Gaussianity of the score, and (C) the part carried by its correlation across intervals; their product reproduces the total to within 1% over the whole grid. Each block is a grid of the design plane in the orientation of Figures 4 and 6: the channel count increases to the right, the dimensionless noise $\tilde{S}$ increases upward, and inside every cell the horizontal axis is $\tilde{\Delta}$ from 0.01 to 1, shared by every panel and labelled once per block. The vertical axis is the distortion factor on a logarithmic scale, shared by all twelve columns so that they can be compared, with the green line at 1 marking a report that is neither over- nor under-confident. Colour is the parameter. Bands are 95% percentile bootstrap intervals over recordings, of median half-width 0.03. The two readings the figure supports, and the limits on them, are in the text.

Figure 6 collects the measurements into one plane, the channel number against the instrumen-

tal noise, at a single acquisition interval $\tilde{\Delta} = 0.1$. Its content is the boundaries: the measured cells are the input and the lines are the output. Four are drawn, and they are two questions asked of two methods. Two ask whether a method's reported error bar is honest, one for the boundary-conditioned likelihood and one for least squares, and are drawn where that method's information distortion crosses 1.15, a 7% error on the reported standard deviation. Two ask whether a parameter is reachable at all, and are drawn where the distortion-corrected standard error crosses a factor of two.

Read upward at a fixed channel count, the recording loses capabilities in a fixed order, which gives five regions, numbered from the bottom. Region 0, below the lowest boundary, is where the Gaussian closure is itself misspecified and the recording is not a macroscopic problem; it is the microscopic regime of the Theory seen on this plane. In region 1 the likelihood delivers the unitary current, the channel number and an honest interval while least squares reports an interval several times too narrow; in region 2 the amplitude pair stops being separable and the rates alone survive; in region 3 least squares is calibrated too and buys the same answer more cheaply; in region 4 nothing useful is left. The lowest boundary is the one with a negative slope, measured at -0.56 in $d \log_{10} \text{noise} / d \log_{10} N_{\text{ch}}$, because instrumental noise makes the emission more Gaussian and so relaxes the occupancy closure while hurting everything else. Regions 1 to 4 all sit inside the Gaussian regime, and what orders them is how much the data carry rather than whether the closures hold. The telegraphic regime is a statement about Δ and so is not a region of this plane; it is reached by moving along the interval axis of Figure 4.

The boundary at which the classical error bar becomes honest is the one this paper predicts and then measures. Fitted in logarithms over four measured crossings in each of two parameter directions, it runs as instrumental noise proportional to channel number, with slopes of 1.01 and 1.03. That is the crossover the mechanism predicts, since the gating variance grows with the channel count while the instrumental variance does not, so the ratio of the two decides whether a window-ignoring fit is honest. Here it is a measured slope rather than an assumed scaling.

The ordering of the boundaries is the result the map exists to state. There is no region of the measured plane in which the amplitude still carries information and the classical error bar is simultaneously trustworthy; the two boundaries are separated by one to three decades of noise everywhere measured. Whenever the fluctuations still say something about the unitary current or the channel number, a classical fit reports a confidence interval that is too narrow. The scope of that statement has to be stated with it: the classical arm here is least squares on the deterministic mean current with the unitary current held fixed, in its four-parameter configuration. It is not the variance-mean regression, this paper does not benchmark non-stationary fluctuation analysis, and nothing here should be read as a measured comparison against it.

The map also fixes where a parameter stops being reachable at any noise level. Minimising the corrected standard error over the whole noise sweep, the channel number is unreachable below about 17 channels and the opening rate below about 52, at this acquisition interval. The unitary current and the closing rate have no such floor within the simulated range.

The evidence is thinner than the picture, and the subsection says so rather than leaving it to a caption. The two boundaries at which least squares becomes honest, and the closure floor, are each fitted through four measured crossings. The channel count at which the third region opens spans 120 to 1368 depending only on whether the strict or the loose criterion is used. The lower vertex is extrapolated to N_{ch} between 3 and 9, below the simulated floor of ten channels, and the ten cells that would pin it have not been run. All the boundaries are drawn at one acquisition interval, with the ribbons showing where each moves as the interval is swept from 0.01 to 1 in units of τ .

Discussion

Over most of the design space the boundary-conditioned member reports an uncertainty close to the one it delivers, and where it does not, the departure is measured rather than assumed. Of the




Figure 6. What a macroscopic recording can be asked for: a usage map over channel number and instrumental noise. The plane is the design space, the number of channels in the patch against the dimensionless instrumental noise, both on logarithmic scales, at one acquisition interval $\tilde{\Delta} = 0.1$. Four boundaries divide it, and they are two questions asked of two methods. Colour is the parameter the boundary is about. Line type is the question: solid asks whether the parameter can be measured at all, drawn where the distortion-corrected standard error crosses a factor of two; dotted asks whether the reported error bar of least squares is honest, drawn where its information distortion crosses 1.15; dashed asks the same of the likelihood's own error bar, and is the boundary below which the Gaussian closure is misspecified. Ribbons are the same boundary swept over $\tilde{\Delta}$ from 0.01 to 1, so a ribbon says where a reader's own sampling rate puts the boundary and is not an error bar. Dots are the measured crossings and the line through them is their fitted power law, drawn faint where it leaves the simulated range, which is marked. Vertical strips mark the two vertices, each a range set by the strict and the loose criterion. The boundaries partition the plane into five regions, filled and keyed by what a recording in each can be asked for, and the three labelled boxes mark where common preparations sit; the regions are named and read in order in the text. The map is a concept map, not a phase diagram. Its boundaries are level sets of continuous diagnostics, so a looser criterion moves each of them by up to a decade in noise without changing the layout.

2268 points at which the boundary-conditioned member was measured, one per parameter per design cell over 560 cells, spanning eleven channel counts from 5 to 10^4 and twelve instrumental noise levels with seven acquisition intervals resolved inside each cell, 2116 (93%) have a reported uncertainty within 15% of the one delivered. An experimenter who does not want to consult a map can use that member everywhere, and where its reported interval is the wrong width is the few-channel corner. Where it is wrong, the correction that repairs it is the one the diagnostic already supplies: at $N_{\text{ch}} = 10$ its nominal 95% region covers 0.914 of the estimates under the covariance it reports and 0.947 to 0.949 under the sandwich, over a thousand fits in six parameters.

That member is an approximation and is not exempt from the failure this paper is about. Over the same grid its distortion runs from 0.645, at $N_{\text{ch}} = 10$ with $\tilde{S} = 0.005$ and $\tilde{\Delta} = 0.01$ on the channel-number direction, where it over-states its own uncertainty, to 1.701, at $N_{\text{ch}} = 5$ with $\tilde{S} = 0.01$ and $\tilde{\Delta} = 1$ on the closing rate, where it under-states it. The departure is two-sided, and that is how it should be read: the reported interval moves in both directions as the channel count falls, and the approximation does not become safe by erring conservatively on average. Both excursions sit where the open population is a small integer count whose distribution is visibly discrete, which is where the occupancy Gaussian, the first of the family's two closures, is expected to degrade. No channel count is named as the point at which that closure becomes misspecified, and the reason is a limit of this design rather than a gap in the measurement: with the open probability held at one half throughout, a channel count and an open-channel count differ by a fixed factor of two that nothing here can separate, so a threshold stated in channels would be stated in the wrong unit. What is measured instead is where each parameter stops being reachable at all, below 17 channels for the channel count and below 52 for the opening rate. The boundary also moves with the instrumental noise, downward as the noise rises, because instrumental variance makes the emission more Gaussian and so relaxes the closure while hurting everything else.

A second boundary of the family's validity is argued here from its premise and has not been measured. Every member above least squares assumes that the instrumental contribution to a recorded interval is Gaussian and white, entering the predictive variance as one additive constant and leaving no memory between intervals. Recordings depart from that description, through the two-level switching of a trapped charge in the headstage and through the flicker component of the amplifier, and the discrepancy literature for models of this class already argues that unmodelled structure of exactly this kind, rather than independent Gaussian error, is what the residuals of a fitted electrophysiology model carry *Lei et al. (2020)*. Where such a departure is large, the diagnostic would report a distortion that belongs to the noise model rather than to the gating closure, and the score alone does not separate the two. We have not run it. The simulator emits Gaussian white instrumental noise only, no cell with any other noise process exists in the sweep, and nothing in this paper measures the size of the effect. This is an expectation from the closure's premise rather than a result, and the way to settle it is to simulate the alternative noise process and repeat the same measurement.

The same amplifier stage carries a second departure of the same shape, and this one we adopt knowingly. A recording is low-pass filtered before it is digitised, normally by a Bessel filter, so the measurement kernel of a real experiment is that filter's impulse response, while the observable used here, in the simulator and in every likelihood alike, is the uniform average over the acquisition window (Theory). The uniform window is the first approximation to that kernel and it is the one that makes the family well posed, since it keeps all of the weight of a recorded value inside a single interval, which is the interval whose two ends the calibrated member conditions on. A Bessel kernel moves both premises of this paragraph at once: it spreads weight across interval boundaries, so successive recorded values are correlated through the signal itself, and it colours η_i . Neither effect is measured here. Extending the recursion to a general kernel with its coloured noise, and then re-running this same diagnostic against it, is the natural next step, and until it is taken the numbers above describe uniformly averaged recordings. That step is not a small one: convolving an amplifier's kernel explicitly, at a 10 kHz Bessel setting with 50 kHz digitisation, was es-

timated to raise the demand of a nine-scheme Bayesian comparison by more than tenfold **Moffatt and Pierdominici-Sottile (2025)**.

What buys the calibration is conditioning at both ends of the acquisition interval. Conditioning on the start alone does not buy it, and at the reference cell it costs. The one-endpoint member and the boundary-conditioned member assign the same total variance to the observed interval average, from the same state, and they book it differently: the spread of the interval mean across end states enters the one-endpoint member as observation noise, while the boundary-conditioned member carries it as a covariance between the interval current and the occupancy at the end of the window, which is the numerator of the gain in Eq. 22. The whole distance between the two therefore lies in the update. On the ellipse-area measure at $N_{\text{ch}} = 100$, $\tilde{S} = 0.01$ and $\tilde{\Delta} = 0.1$, the one-endpoint member reads 1.63 against 1.01 for the boundary-conditioned one, and it is also worse than the instantaneous recursive filter at 1.17, so the intermediate rung does not interpolate between the two ends of the step it sits in.

A supplement takes that step in its two algebraic halves and confirms a prediction that was registered before the run. The variance-recursive member removes the end-state spread from the observation variance without restoring it in the gain, and the prediction was that its reported uncertainty would come out worse than the one-endpoint member's, over-confident and more so. It does. The reading carries one qualification the algebra makes obligatory: the predictive variance divides the gain, so altering the variance form also moves the update, and the two halves are not a clean separation into a variance step and a gain step. This is a supplement result and the body's claim does not rest on it.

The diagnosis is Milesescu's, now measured. **Milesescu et al. (2005)** named the local time correlation of the current as the dominant error of estimates that do not filter, and that correlation is what the distortion decomposition isolates. Interval by interval every rung reports its information correctly; the miscalibration appears in how the intervals accumulate. It is visible directly in the autocorrelation of the score, which at lag one on the closing rate is -0.0043 ± 0.0043 for IR, indistinguishable from zero, against 0.191 ± 0.004 for R, 0.864 ± 0.0019 for NR and 0.856 ± 0.0022 for either least-squares arm (six-parameter configuration, Methods). Conditioning on both interval ends leaves a white score. A likelihood member that skips an endpoint reads sequential structure as independent evidence, which is how a per-interval identity that holds everywhere can accumulate into an information ratio off by more than an order of magnitude. The accumulated ratio of least squares is not read this way: its information presumes one constant residual variance, which the simulator violates, so its factor of fifteen measures the size of that break rather than a failure of accumulation, and it sits above the calibrated value at the per-step level too (Diagnostics).

Del Core and Mirams (2025) argue against filtering on cost, since the moment integration and the correction step run between every pair of samples. The cost is real, and what it buys is now quantified in units of misreported uncertainty. How much cost it is depends on the scheme and the implementation rather than on anything measured here, and one published figure gives its scale: a Bayesian comparison of nine schemes for a ligand-gated receptor with a member of this family took twelve to fifty-two days per scheme on sixteen to thirty-two cores **Moffatt and Pierdominici-Sottile (2025)**. Where the instrumental noise is large next to the gating variance the cheaper members recover: at $N_{\text{ch}} = 10^4$ the closing-rate distortion of the instantaneous recursive filter falls from 1.42 at $\tilde{S} = 0.01$ to 1.03 at $\tilde{S} = 10$ and 1.00 at $\tilde{S} = 100$, and least squares follows the same path about three decades higher in noise. In that part of the plane the filter's expense is hard to justify on calibration alone. Toward few channels and short intervals the correlation part of the distortion grows, and the degrees-of-freedom count a non-conditioned fit uses is wrong by a factor that no compute saving offsets: at $\tilde{S} = 0.01$ the integrated autocorrelation of the standardized residual is 9.26 for least squares and 3.93 for the open-loop member, against 1.008 for the boundary-conditioned filter, so the classical fit behaves as though it had about nine times more independent observations than the recording supplies. The first moment is a separate charge on the same trade. The instantaneous filter carries a bias in the channel number of 35 to 40% that does not shrink over four decades of

channel count, and no error-bar correction touches it. The map states both halves, so a reader can locate a recording on it and see which side of the trade it sits on.

One result speaks to how a recording should be spent. Once the agonist is removed and the open population relaxes, a member that has conditioned on the past gains nothing further about the channel number or the opening rate, whose information reaches the numerical floor within six intervals of the washout, while the closing rate and the unitary current stay informative through the decay by different routes, the first through its shape and the second through the fluctuation of the channels still open (Figure 3–figure supplements 1 and 2). The shape does not hold for every rung, and that is itself the finding: the non-recursive members keep about a sixth of their pulse information about the channel number through the washout against under one per cent for the recursive ones (Results), so this is a statement about conditioning and not about the observable. It answers a question *Del Core and Mirams (2025)* leave open. For the bench it says that the transient carries what the decay cannot: extending a record after the agonist is removed adds data without adding information about N_{ch} or k_{on} , so an experiment that wants those two should spend its samples on the jump and its rise, while an experiment that wants the closing rate or the unitary current is served by the decay as well.

The ordering of the boundaries on the map is what an experimenter can act on. There is no region of the measured plane in which the fluctuations still carry information about the unitary current or the channel number and the classical error bar is simultaneously trustworthy; the two boundaries are separated by one to three decades of noise everywhere measured, and the boundary at which the classical interval becomes honest runs with fitted log-log slopes of 1.01 and 1.03, so it scales as instrumental noise proportional to channel number. Each of those boundaries is fitted through four measured crossings at one acquisition interval. The map is a concept map, not a phase diagram. Its boundaries are level sets of continuous diagnostics, so a looser criterion moves each of them by up to a decade in noise without changing the layout. That scaling was predicted from the mechanism, since the gating variance grows with the channel count while the instrumental variance does not, and it is then measured as a slope. The concession that travels with it is that least squares recovers the rates, with a rate bias at or near zero over the same plane; what it gets wrong is how well it says it recovered them. The scope of the comparison travels with it too. The classical arm of the design-plane and map results is ILSE in the four-parameter configuration of Methods, least squares on the interval-averaged deterministic mean current with the unitary current held fixed; the time-resolved comparison uses the six-parameter configuration and carries both least-squares arms, and no least-squares quantity is carried across the two configurations. This paper does not benchmark non-stationary fluctuation analysis, no variance-mean regression is run anywhere in it, and nothing above is a measured comparison against that method.

As an algorithm, MacroIR is an integrated-measurement augmented Kalman filter, a device long known in control and target tracking *Zadrozny (1988)*, and its predictions coincide with that construction to about 10^{-8} . What is specific here is the realization for an exact continuous-time Markov chain of channel occupancies, its use at the populations simulated in this paper, from 5 to 10^4 channels, where the single-molecule integrated filters do not reach, and the measurement of what the cheaper members cost. The device is conceded; the continuous-time realization derived here and the validity map are the contribution, the recursive lineage they extend is *Moffatt (2007)* and *Münch et al. (2022)*, and the family's first application to macroscopic recordings is *Moffatt and Pierdominici-Sottile (2025)*.

Münch et al. (2022) asked this paper's question of a filter of this family, and answered it by counting: over repeated simulated data sets they recorded how often the true rate matrix falls inside a highest-density credibility volume of a given probability mass, and compared that count against the binomial rate an exact posterior would produce. Their filter holds the rate across masses from zero to 0.95 where the rate-equation approach does not, and they proposed the exercise as a general way of assessing an algorithm's validity. Three things separate it from what is done here. What their count tests is a posterior, so the prior and the sampler are inside the verdict together with

the likelihood, and their own microscopic-reversibility prior is visible in the behaviour of one equilibrium constant; the identities used here are evaluated on the likelihood alone, at parameters that are known, with no prior and no sampler. What their count returns is a probability that a region holds the truth; what is returned here is a matrix, which says by how much and in which parameter the reported uncertainty departs and separates what each interval commits from what accumulates across them. And a coverage rate costs one posterior fit per replicate at every design point, which is what keeps such a study to a line of conditions, while a score and an information evaluated at known parameters cost one pass each, which is what puts 560 cells within reach. On the two axes of this study their filter maps the instantaneous occupancy to the observation rather than conditioning on the ends of an acquisition interval, which places it with the instantaneous recursive member; what it generalizes is the emission, through state-dependent open-channel noise and a second, fluorescence observable, and it is that generalization that costs them the analytical update and obliges a closure at two moments of their own. The window they treat in the other direction, recommending minimal analogue filtering and a higher sampling rate in place of a model of the acquisition kernel, which moves an experiment along the interval axis swept here.

There is a reason this matters more as the inference improves rather than less. A method that cannot resolve a mechanism cannot mislead anyone about one. When we proposed that $P2X_2$ passes through a flip state before opening, the claim carried a signature that can be pointed to in the record, a delay before current appears, and other groups later confirmed it *Moffatt and Hume (2007)*. The asymmetric activation reported from the same recordings two decades later has no such signature and rests on ratios of evidence between schemes *Moffatt and Pierdominici-Sottile (2025)*. A conclusion of that kind is a function of the likelihood that produced it, so whatever an approximation gets wrong about the information in a record passes into the mechanism it selects, and the resolution that makes such a conclusion reachable is what makes its calibration worth measuring.

The full distortion matrix is reported, rather than a scalar summary, because it is the input to the next question. Near the maximum the same matrix propagates into the Bayesian evidence through a volume term $\frac{1}{2} \log \det \mathbf{C}$ and an effective-sample rescaling $\alpha^* = p / \text{tr } \mathbf{C}$, so whether model comparison under an approximate likelihood ranks mechanisms correctly is decided by the object measured here. That correction is derived elsewhere and is stated only as motivation. It is one bridge of a programme that assembles exact simulation, an approximate likelihood and its correction into inference on real recordings *Moffatt and Pierdominici-Sottile (2025)*. The two orderings of the $P2X_2$ comparison cited above were produced by likelihoods that differ in the property measured here, and what this paper supplies is how far the reported information of such a likelihood sits from the information it delivers. Which of those two orderings is right is a further question: it needs the volume and effective-sample terms above and the evidences recomputed, and neither is done here.

A reader who accepts the miscalibration can still ask what it would cost to repair it without changing likelihood, and the alternatives answer different questions. The standard error across independent fits to separate patches is free, since the fits already exist, but it describes the uncertainty of a population mean across patches, mixes statistical with biological variability, and leaves any bias in place while shrinking as the square root of the number of patches. Resampling whole sweeps and refitting is valid and correctly preserves the within-sweep correlation, and it gives an honest interval around whatever the estimator converges to while saying nothing about whether that estimator is the right one. Resampling residuals is not valid here, because independent resampling destroys the correlation that is the problem. A parametric bootstrap assumes the model one was trying to check. A block bootstrap needs a block length well above the correlation time, so it needs that correlation time measured before it can be configured. Simulating from the fitted model and forming the sandwich answers the same question the diagnostic does, and its cost is set by the precision needed on a covariance rather than by the model, which puts it orders of magnitude above a single evaluation and still leaves it simulating from the model one wanted to

verify.

Where exchangeable sweeps are plentiful, the classical route is strong and should be said to be. Non-stationary fluctuation analysis returns the channel number and the unitary current with no likelihood at all, and a sweep-level bootstrap returns honest intervals with no model, and between them they cover much of routine electrophysiology. The case for a filter is strongest where that route is weakest: few sweeps, rundown, conditions that are not exchangeable, a single precious recording, and an experiment that wants kinetics and amplitudes jointly from the same record.

What this paper does not settle

The broad question, when to prefer one of these methods over another, cannot be closed with the elements assembled here, and the honest move is to say which parts of it are missing rather than to answer it badly. Four of them depend on facts about the reader's experiment that no simulation study supplies. The cost comparison in general depends on the number of states, the implementation and the optimiser, and this paper has one model, one implementation and one machine. Whether a sweep-level bootstrap suffices depends on how many sweeps a preparation tolerates and whether they are exchangeable, which is a rundown question. How much precision is lost by fitting with least squares is measured here for this model, and whether that loss matters is the reader's judgement about their own question. How any of this composes with non-stationary fluctuation analysis as a complete workflow is not addressed.

The fifth is about the diagnostic itself and bounds what it verifies. The comparison here is between two independent implementations, the likelihood and the exact simulator, and everything downstream of the point at which they separate is tested by it: the propagation, the update, the emission arithmetic, the accumulation. Upstream of that point they share one specification, the rate matrix built from the parameters, the conductance vector, the protocol and the sampling times, and a diagnostic built on their agreement cannot see an error there. For a two-state scheme that specification is small enough to be read and checked by eye, and it is written out in Methods for exactly that reason. The general remedy is different in kind: verify each object against the theorems it must satisfy, the semigroup and Chapman-Kolmogorov identities for the propagator, the stationarity condition, the normalisation of the occupancy, and the moment identities of the emission. That is a piece of work with its own design and it is not undertaken here. The sixth is the scope declared in the paragraph below.

What this changes is what the paper hands over. The reader supplies their sweep count, their rundown, their model size and their tolerance, and decides; what they cannot measure for themselves, and what is supplied here, is whether a reported interval is the interval that is delivered.

What produces the boundary-conditioned member's own departure is not known, and the shape of the ignorance is worth stating because it is sharper than a general disclaimer. The decomposition puts the departure in the term that accumulates across intervals rather than in the per-sample one: the member's per-sample non-Gaussianity stays close to one where its total does not. Five candidate criteria for where the closure should fail were derived and measured against the observed boundary, and every one of them was a criterion on the per-sample cumulants of the occupancy, so every one of them failed to locate it. What survives from that exercise is a bound and not a position: occupancy non-normality alone cannot produce a factor-two distortion anywhere on the measured grid. The measured channel-count exponent of the boundary is -0.25 , which matches neither the -0.5 a skewness mechanism would give nor the -1.0 of a kurtosis one, and it is not a clean power law. Stated as a conjecture and in those words: the limits behave like a balance that has to come out just right, several quantities moving together in a way that leaves the departure small over four decades, which is the signature of a conserved quantity or a comparable invariant that has not been identified here.

The perimeter of these claims is a model, and it is declared as one. The scheme is two states with the open probability held at 0.5; the protocol is a single concentration jump; the data are simulated, because a diagnostic that compares what a likelihood delivers against what it reports

needs a known truth to mean anything; and the inference is likelihood only, with no prior and no evidence computation. The two-state scheme isolates the closure's own error with no confounding from scheme misspecification, and the fixed open probability means that a channel count and an opening count differ throughout by a factor of two that no measurement in this paper can settle. What the channel-number axis carries at $p = 0.5$ is the gating variance, so a reader whose receptor sits at another open probability can match it by entering the map at $4Np(1 - p)$ in place of the channel count N ; the substitution is exact for the variance and says nothing about the skewness of the open population, which vanishes at $p = 0.5$ and which no cell here probes. Richer schemes, the stationary regime, the few-channel boundary, the filtered measurement kernel and experimental data are named companions rather than omissions. Two of those companions a reader need not wait for. The software carries the calibrated member with its score and its Fisher information and the exact simulator, so the two identities can be run on another scheme, at another open probability, by whoever wants the answer for their own model rather than an extrapolation from this one; what it does not carry is least squares, the open-loop member or the recursive member on instantaneous samples, so the comparison between rungs is reproducible only from the engine.

Materials and methods

Minimal model and simulation

We work with a two-state channel scheme, closed (C) and open (O), the smallest scheme in which the acquisition-averaging effect we study is present. The closed-to-open transition is agonist dependent with association constant k_{on} multiplying the agonist concentration; the open-to-closed transition is agonist independent with dissociation rate k_{off} . Only the open state conducts, and the code applies a sign flip so the single-channel current is inward. At the start of every recording all channels are closed, so the occupancy is $(1, 0)$.

The scheme carries six parameters, all estimated in base-ten logarithmic coordinates. The values used to generate the data reported here, read from the production run scripts rather than from the parameter tables in the repository (the tables hold stale values the figures never used), are $k_{\text{on}} = 10 \mu\text{M}^{-1} \text{s}^{-1}$, $k_{\text{off}} = 100 \text{s}^{-1}$, a unitary conductance g set to 1 (the engine parameter `unitary_current`, written i where the noise axis is defined below), and a current baseline b set to 1 (both in units of g), while the white-noise level and the mean channel number were swept (below). Current is expressed throughout in units of the unitary conductance g , which corresponds physically to a single-channel current of roughly 1 pA. The dissociation time $\tau = 1/k_{\text{off}} = 10 \text{ms}$ sets the natural time scale that normalises the acquisition-interval axis. The relaxation at the agonist condition is faster, with rate $k_{\text{on}}[A] + k_{\text{off}} = 200 \text{s}^{-1}$; the axis uses $\tau = 1/k_{\text{off}}$ by convention.

The emission model is white instrumental noise only. Both the pink (one-over- f) and the proportional noise channels are set to zero, so the term the reader may expect from flicker noise is genuinely absent. For a measurement interval t of duration Δ_t , the predicted current has mean $\bar{y}_t = N_{\text{ch}} \mu_t \gamma_t + b$ and variance

$$\sigma_t^2 = e_t + N_{\text{ch}} v_t, \quad e_t = \text{Current_Noise} \frac{f_s}{n_{\text{samp},t}} = \frac{\text{Current_Noise}}{\Delta_t},$$

where N_{ch} is the channel number, μ_t the occupancy mean, γ_t the interval-averaged conductance, f_s the raw sampling rate, $n_{\text{samp},t}$ the number of raw samples averaged into the interval, e_t the white instrumental variance over the interval, and $N_{\text{ch}} v_t$ the channel-gating variance. The white variance e_t falls as the interval lengthens, because more raw samples are averaged into each recorded value. The gating variance v_t is the single-interval conductance variance, whose form differs across the algorithm family and is derived in the Theory section.

The stimulus is a single concentration jump in three segments, a pre-pulse at $0 \mu\text{M}$, an agonist step at $10 \mu\text{M}$, and a post-pulse at $0 \mu\text{M}$, with raw sampling at $f_s = 50 \text{kHz}$ throughout. Each recorded measurement is the mean of $n_{\text{samp},t}$ consecutive raw samples, so the averaging window is fixed at acquisition. That window is the acquisition interval the whole study is about, and its length

is set by the number of raw samples entering each recorded value. The averaging is uniform over the window and no analogue filter stage is simulated, so a recorded value is a boxcar average of the trajectory, which is the observable every likelihood in the family is written for (Theory). Recordings from a real amplifier are Bessel filtered instead, and that kernel is outside the scope of this paper (Discussion).

Ground truth is generated by exact simulation of the continuous-time Markov chain (CTMC) by uniformization, so the interval average of the conductance is realised from an actual sampled trajectory rather than from any moment-closure approximation. Uniformization draws the transition times themselves, so there is no within-interval discretisation to choose and none to justify. The production scripts do pass a `number_of_substeps` argument, and on this branch it has no effect: the simulation parameters are built from the algorithm name alone, and the substep count they carry is zero. This is what licenses treating the simulation as the reference against which the Gaussian likelihood approximations are judged: the approximations are the object under test, and the simulator supplies the exact answer they are compared to. Every maximum-likelihood grid cell reported here, the least-squares arm included, used 10^4 recordings, with one exception: the cell that carries the coverage result ($N_{\text{ch}} = 10$, $\tilde{S} = 0.005$, boundary-conditioned member) was run at 10^5 , because coverage of a nominal 95% region in six parameters needs a thousand independent fits and those fits are taken at group size 100, so 10^5 recordings are required to produce them. The time-resolved run used 1,000 recordings; the single-recording illustration (Figure 1) was drawn by the substep sampler instead, one 20-channel trace at 250 substeps per interval, and does not enter the quantitative comparison.

The simulation seed was set to 0, which in this engine is the sentinel for a random seed drawn from `std::random_device`, and the resolved seed was never logged. The deposited ensembles are therefore statistically equivalent to, but not bit-for-bit reproducible from, a rerun of the same script. Because each ensemble is large (1,000 to 10,000 recordings), the reported statistics are stable; the Monte-Carlo standard here is statistical equivalence, not bit identity.

The likelihood family and its members

The members compared here answer one question first, then a second. The first question is whether the gating fluctuations are modelled at all. Classical nonlinear least squares answers no: it fits the deterministic mean current and treats every departure from it as independent measurement error. The macroscopic likelihoods answer yes, and among them the second question is how much of the acquisition interval the interval-averaged conductance is conditioned on. Least squares is the bottom rung of that cost ladder and the comparison anchor of this paper.

Least squares is selected by the family flag *family approximation* = 2 and runs non-recursive; the variance flag is accepted and ignored on this branch. The averaging flag still applies to it, and it is what separates the two least-squares arms: at 0 the fit predicts the deterministic mean current at one instant, the middle of the acquisition window (LSE, data key `nonlinearsqr_g`), and at 1 it predicts that same mean averaged over the window (ILSE, data key `nonlinearsqr`). shares its mean model exactly with the macroscopic member carrying the same averaging flag, LSE with NR and ILSE with INR, and differs from it only in the predictive variance: on the recording of Figure 1 the predicted means of LSE and NR agree to 9×10^{-16} interval by interval, and those of ILSE and INR to 2×10^{-16} , while the two least-squares arms differ from each other by 5×10^{-2} on the same intervals. Both arms sit off the lattice of endpoint and recursion flags that organises the macroscopic members; the I prefix records the window setting and does not place ILSE on that lattice.

Every cell of the design grid is run with ILSE: the un-averaged arm was never dispatched on the grid, and there is no `nonlinearsqr_g` output under `figures/data`. LSE is measured at the single cell of the time-resolved figure, from the same script that produces the ensemble the other members score, and drawn on the illustrative recording of Figure 1. The macroscopic members all take *family approximation* = 0 and are one filter run under flags that set how the interval-averaged conductance is conditioned. The *recursive* flag sets whether the occupancy covariance is propagated

between intervals. The *averaging* flag $av \in \{0, 1, 2\}$ counts the number of interval endpoints the hidden state is conditioned on, and it is the program's name for the window axis of the Theory section: 0 an instant, 1 the state at the window's start, 2 the states at both of its ends. A third flag, the *variance form*, distinguishes the two members that condition on the interval start: MR predicts the interval variance with the total per-start-state conductance variance, while VR removes the between-endpoint part, keeping MR's mean and gain.

Member	Data key	Recursive	av	Variance	Conditioning of the conductance
LSE	nonlinearsqr_g	no	0	n/a	none; the deterministic mean current at the sampling instant
ILSE	nonlinearsqr	no	1	n/a	none; the deterministic mean current averaged over the interval
NR	macro_NR	no	0	n/a	instantaneous; the acquisition averaging is ignored
INR	macro_INR	no	1	total	interval mean, without recursion; the endpoints are not separable
R	macro_R	yes	0	n/a	instantaneous; the acquisition averaging is ignored
MR	macro_MR	yes	1	total	interval mean given the start state
VR	macro_VR	yes	1	residual	MR's mean and gain, with the between-endpoint variance removed
IR	macro_IR	yes	2	residual	interval mean given both boundary states; the endpoints enter the gain

Table 2. The members as flags. Both least-squares arms take *family approximation* = 2 and every macroscopic member takes *family approximation* = 0. Reading the endpoint flag as a ladder gives no endpoints (NR, R), one endpoint (MR, VR), then both boundary states (IR); INR sits outside that ordering, because without a recursive update the boundary state is unobservable and conditioning on the start state and conditioning on both ends describe the same method. The variance column reports the form in force rather than the dispatched flag: IR is dispatched with the total flag, but with both endpoints conditioned on the residual form is used regardless of it, because the between-endpoint spread is already carried by the boundary term. Even so, taken from the same state MR and IR assign the interval average the same total variance, by two decompositions that rearrange into each other (Appendix 1), so what separates them is the gain; because that variance divides the gain, the states they carry part company at the first gated update and the variances they report along a recording part with them. VR keeps MR's mean and boundary-free gain and replaces the total per-start-state conductance variance with the residual one, which is not the same thing as carrying IR's variance. written by different producers, so their output files carry different prefixes, *figure_3_LSE_* and *figure_3_G_*. INR is the member published as MacroINR in *Moffatt and Pierdominici-Sottile (2025)*, where it served as the control against which the boundary-conditioned member was compared; it is reported here only in the supplement.

Table 2 lists the members. Reading the endpoint flag as a ladder, the macroscopic members condition the conductance on progressively more of the interval: no endpoints, one endpoint, then both boundary states, with the exact trajectory (what the simulator supplies) as the unreachable top rung. IR is the top implemented rung: it conditions on the pair of channel states at the two ends of the interval, which we call the boundary state, and marginalises the interior analytically. It is specified by the construction of the Theory section and is placed here as one member of the completed grid; an interval-recursive likelihood of this family was first applied to macroscopic recordings, of the P2X₂ receptor (*Moffatt and Pierdominici-Sottile, 2025*). It coincides with a time-integrated Kalman filter, a correspondence stated and quantified in the Discussion. The construction of the interval-averaged mean and variance, the boundary-state covariance, and the conditioning step that consumes them are derived in the Theory section and its supplement; we do not restate them here, beyond the regularizer of Eq. 40 below.

Two reductions are worth stating because they anchor the family to prior work. Setting the boundary-conditioned conductance variance to zero in the recursive, no-endpoint member (R) recovers the macroscopic-fluctuation update of *Moffatt (2007)*, so R is the recursive ancestor already in print. The single structural difference between the M and I families is that the one-endpoint members drop the boundary cross-covariance term $N_{ch} \mathbf{y}^T \mathbf{\Sigma} \mathbf{y}$ that IR retains, the tilde being the boundary-pair contraction defined at Eq. 20. Every member is written out on one template in Appendix 1.

Throughout the reported runs the remaining build flags were held fixed: *taylor variance correction* off, *micro approximation* off, and the *variance approximation* on. The V in VR names the *variance form* flag, set to 1, and has no connection to the *taylor variance correction* flag, which is off in every run reported here; the Taylor-corrected variants MRT and IRT are not part of this paper. Because the Taylor variance correction is off, the recursive one-endpoint member that runs is MacroMR and not MacroMRT.

The filter carries three numerical safeguards, present because the code runs them and because anyone reimplementing it meets the same constraints on day one. None is housekeeping: the first fixes what the boundary-conditioned moments are where a transition cannot occur, and without the other two the recursive members return no value at all at the smaller channel counts.

The first sits upstream of the filter cycle, in the construction of the boundary-conditioned moments of Eq. 4. Both are quotients by the transition probability $P_{i \rightarrow j}(\Delta)$, and that denominator is zero wherever the transition cannot happen within one window, which on this design is not an edge case: at zero agonist $k_{\text{on}}[A]$ vanishes identically and $P_{C \rightarrow O}$ is zero over the whole pre-pulse segment. The numerator vanishes with it, so the quotient is 0/0 in exact arithmetic and, in floating point, a ratio of two quantities no larger than the rounding error of the reconstruction that produced them. The implementation does not test for that case. It reads each quotient as a posterior mean and adds a conjugate pseudo-count ϵ to numerator and denominator alike,

$$\bar{y}_{i \rightarrow j} = \frac{[\mathbf{F}_1(\Delta)]_{i \rightarrow j} / \Delta + \epsilon \frac{1}{2}(\gamma_i + \gamma_j)}{P_{i \rightarrow j}(\Delta) + \epsilon}, \quad \bar{m}_{i \rightarrow j} = \frac{2[\mathbf{F}_2(\Delta)]_{i \rightarrow j} / \Delta^2 + \epsilon \frac{1}{2}(\gamma_i^2 + \gamma_j^2)}{P_{i \rightarrow j}(\Delta) + \epsilon}, \quad (40)$$

where $\bar{m}_{i \rightarrow j}$ is the second moment of the interval-averaged conductance over the same boundary pair, and the variance follows by subtraction, $\bar{v}_{i \rightarrow j} = \bar{m}_{i \rightarrow j} - \bar{y}_{i \rightarrow j}^2$. The two prior means are the end-point averages of Eq. 9, so the prior the variance inherits is $((\gamma_i - \gamma_j)/2)^2$ and is never imposed separately. The weight is $\epsilon = \epsilon_{\text{mach}} \kappa_F(\mathbf{V})$, machine epsilon times the Frobenius condition number of the matrix of right eigenvectors of $\mathbf{Q}$, which bounds the floating-point noise the spectral reconstruction can leave in $\mathbf{P}(\Delta)$. The pseudo-count is therefore inert wherever the transition probability stands above that noise, takes over where the quotient carries no information, and introduces no fitted constant. Both choices in it are deliberate. Keying ϵ to $\kappa_F(\mathbf{V})$ ties the regularizer to the conditioning of the decomposition that produced the noise rather than to any model constant, and taking the prior from the conductance vector rather than from the marginals of $\mathbf{F}_1$ keeps it clean when $\mathbf{V}$ is ill conditioned, which is when those marginals carry the same reconstruction error the pseudo-count exists to absorb and when the prior is most needed.

The second holds the occupancy mean on the probability simplex. Nothing in the Gaussian conditioning of Eq. 22 knows that $\boldsymbol{\mu}$ is a probability vector, and an undamped step can drive its entries negative or above one, so the mean step is scaled by a trust coefficient $\alpha_\mu \in (0, 1]$, the largest scaling that keeps the posterior on the simplex. The covariance is treated separately and takes its full step, since the rank-one down-date of Eq. 23 is positive semidefinite by construction; the two trusts are not tied to a common factor. The form of the damping matters here in a way it would not in an ordinary filter. This paper differentiates the log-likelihood with respect to $\boldsymbol{\theta}$, so a coefficient that is merely continuous is not enough: a hard $\alpha = \min(1, c \alpha_p)$ leaves a step in $\partial \alpha / \partial \boldsymbol{\theta}$ at the boundary and pumps variance into the score whenever the system sits near it, and a discontinuous form leaves a delta. The coefficient is therefore taken in the smooth form $\alpha = \frac{1}{2}(1 + c \alpha_p - \sqrt{(1 - c \alpha_p)^2 + \epsilon^2})$, a softened minimum of 1 and $c \alpha_p$ that is infinitely differentiable in $\boldsymbol{\theta}$ everywhere. Here α_p is the largest scaling that keeps the posterior mean on the simplex, $c = 0.9$ is a safety factor applied when the constraint binds, leaving a ten per cent margin against the boundary, and $\epsilon = 10^{-4}$ sets the width of the smoothing band, so the residual bias of α below the hard minimum is at most $\epsilon/2$, which is three orders of magnitude smaller than that margin. Reimplementing this with a hard minimum reproduces the mean current and corrupts every quantity this paper measures.

The third is a pair of constants: a binomial-count floor of 5, below which the expected count in a state is too small for the Gaussian treatment of that state and the member falls back to its

non-recursive form, and a minimum occupancy probability of 10^{-12} .

The two least-squares parameter configurations

Least squares runs in two parameter configurations, and they are different statistical models. Both select *family approximation* = 2, both fit data from the same simulator, and both minimise the same sum of squares. They differ in which parameters are free, so no quantity computed under one may be compared against the same quantity computed under the other. Which parameters are free is independent of the window setting of the previous section, so the configurations and the arms cross: the six-parameter configuration is run on both arms and supplies the least-squares pair of Figure 3, and the four-parameter configuration is run on ILSE alone and supplies every grid cell.

The residual scale is not a free parameter of either and is not given a value. It is integrated out under a scale-invariant prior, so with $SSE = \sum_i (y_i - \mu_i(\theta))^2$ over the n finite recorded values the objective is the marginal

$$\log L(\theta) = \log \Gamma\left(\frac{n}{2}\right) - \frac{n}{2} \log \pi - \frac{n}{2} \log SSE(\theta), \quad (41)$$

whose maximiser is the ordinary least-squares one, since only the last term depends on θ . What the marginalisation changes is everything downstream of the point estimate. The score is the derivative of Eq. 41 and carries the factor n/SSE rather than a $1/\sigma^2$ that would have to be supplied, and the information the fit reports is

$$\mathbf{H}_{LSE} = \frac{n}{SSE} \sum_i \frac{\partial \mu_i}{\partial \theta} \frac{\partial \mu_i}{\partial \theta}^\top, \quad (42)$$

the Gauss-Newton curvature with $\hat{\sigma}^2 = SSE/n$ entering as a plug-in prefactor and not as a coordinate. There is therefore no $\partial \sigma^2 / \partial \theta$ term to keep or drop: the scale direction is not in the parameter vector at all, which is also why the classical arm carries no score in the noise direction and why its mean squared standardized residual is pinned at one by construction.

The six-parameter configuration feeds the time-resolved figure (Figure 3). Both arms of it are produced by `ops/local/figure_3_time.macroir`, the same script that simulates the ensemble and scores the macroscopic members, so all eight members of that figure score one common set of recordings. All six parameters are free in base-ten logarithmic coordinates: k_{on} , k_{off} , the unitary current, `Current_Noise`, the current baseline and the mean channel number. The script's own header records why. The time-resolved figures accumulate the per-interval Fisher information and never invert it, so the flat directions of the least-squares fit are harmless there; and keeping six free preserves the parameter-index alignment with the macroscopic dumps that the figure code joins against, whereas fixing two would emit four indices against the macroscopic files' six and misalign them without any visible error. Two flat directions are then visible rather than hidden. `Current_Noise` has an identically zero score row and an identically zero Fisher row, because the least-squares mean carries no noise term; and the unitary current and the mean channel number form a rank-one degenerate block along the amplitude ridge $N_{ch} i$, so the unitary current is free in this configuration without being separately identified.

The four-parameter configuration is `ops/local/figure_4_LSE.macroir`, which produces the grid cells, together with `ops/local/figure_3_mle_LSE.macroir`, which is the same model run with the per-group estimate cloud retained for the few cells that need one. Both fix the unitary current at 1 and `Current_Noise` at the cell value, leaving four free parameters: k_{on} , k_{off} , the current baseline and the mean channel number. Fixing those two is forced rather than chosen: they are the flat directions of the previous paragraph, and with them free the parameter Fisher information is singular, so no maximum-likelihood covariance exists. Both fixed values are the values the recordings were simulated from. The script builds one parameter vector and passes it both to the simulator and to the fit as its reference point, so the fixed value and the truth are the same number by construction. Least squares is therefore not handicapped in this configuration but helped: two of the six parameters are supplied to it exactly right, while the likelihood members estimate all six. Where it

is nonetheless the arm whose reported uncertainty is furthest from the delivered one, that failure is not a lack of information about those two. This is the configuration behind every ILSE cell of Figures 4 and 6. Its parameter indices are $0 = k_{\text{on}}$, $1 = k_{\text{off}}$, $2 = \text{current baseline}$, $3 = \text{mean channel number}$. There is no unitary-current index in these files, and index 2 is the baseline.

The consequence is carried into the Results. A score, a Fisher information, a distortion matrix or a log-likelihood computed under six free parameters is a different object from the one computed under four, and the two distortion matrices differ in dimension, 6×6 against 4×4 , so any scalar that accumulates over eigenvalues is not comparable between them. Every statement about least squares in the Results is scoped to one configuration and names it. The statement that least squares cannot deliver the unitary current is a statement about the four-parameter configuration.

Parameter estimation and the two anchors

Inference in each grid cell is two-stage. First, the ensemble is partitioned into groups of consecutive recordings and a separate maximum-likelihood estimate (MLE) is computed per group, producing a cloud of estimates whose scatter is the empirical parameter covariance. The partition is by position and any remainder is dropped, which costs nothing here because the recordings are independent and identically distributed and 10^4 divides by both group sizes used. The optimiser is Gauss-Newton with Levenberg damping, warm-started at the simulation truth. The script sets the damping schedule (initial damping 10, growth factor 3, maximum damping 10^{10}) and three limits: at most 100 iterations, a relative gradient tolerance of 10^{-6} and a value tolerance of 10^{-10} . What stops most fits, however, is a fourth criterion the script cannot reach. The optimiser tests the Newton decrement $\frac{1}{2} g^T H^{-1} g$ first, requiring it below 10^{-8} for the first ten iterations and below 10^{-4} after that; the relaxation is there because past that point a flat likelihood is being chased below its own noise floor. The three script-set limits act as fallbacks behind it, and the gradient-norm one is taken relative to the norm of the warm start rather than of the current θ , so that a parameter running away along a flat direction cannot inflate its own tolerance and be recorded as converged. The Newton-decrement constants are compiled in. Two group sizes were used, 10 and 100; with 10,000 recordings these give 1,000 groups of 10 and 100 groups of 100. Second, a single joint MLE is computed over all replicates at once (one group of size n_{sim}), giving the pooled estimate θ_{pool} ; its near-zero gradient certifies that it is a stationary point of the pooled likelihood.

A group whose refit fails outright is dropped from the cloud, so a cloud can hold fewer estimates than the partition has groups, and its size is conditional on convergence. Every fit that did complete carries its convergence status, decided by the criteria above, into the per-group run log. The figures in the body do not filter on that status; the supplements that restrict themselves to converged fits say so.

Warm-starting the per-group fits at the truth is deliberate. It isolates the error of the likelihood from the error of the optimiser, which is what this paper studies; it would be the wrong choice for a study of an estimator, which this is not.

The two anchors, the simulation truth θ_{sim} and the pooled fit θ_{pool} , are distinct points because the Gaussian macroscopic likelihood is misspecified with respect to the exact simulator. Their difference $\theta_{\text{pool}} - \theta_{\text{sim}}$ is the misspecification bias. Every diagnostic below is computed at both anchors, and the two answer different questions: at θ_{pool} the score is zero by construction so bias washes out, whereas at θ_{sim} the bias is visible but the Fisher information can be indefinite away from the maximum. Each figure states which anchor it uses.

Diagnostics and the Gaussian-Fisher anchor

Faithfulness of each approximate likelihood to the simulator is measured with likelihood-level diagnostics: standardized residuals, the score (the gradient of the log-likelihood), and an information-distortion matrix that compares the covariance J of the score across replicates with the model's own Fisher information. The definitions, sign conventions, and thresholds of these diagnostics are

owned by the Diagnostics section and its metric registry; here we describe only how they were computed.

Every diagnostic reported in this paper is anchored on the analytic Gaussian Fisher information $\bar{G} = \sum_i \text{GFI}_i$, summed over intervals. This is the anchor the Diagnostics section writes $\mathbf{H}$, that is $\mathbf{H} = \bar{G} = \sum_i \text{GFI}_i$. Its per-interval term is

$$\text{GFI}_i = \frac{(\partial_{\theta} \bar{y}_i)(\partial_{\theta} \bar{y}_i)^T}{\sigma_i^2} + \frac{(\partial_{\theta} \sigma_i^2)(\partial_{\theta} \sigma_i^2)^T}{2 \sigma_i^4},$$

which follows from the analytic derivative alone and needs no extra likelihood passes. It is positive semidefinite by construction, so no anchor used here can be indefinite.

A second Fisher construction exists in the codebase and was used in an earlier battery: a central difference of the analytic score, with per-coordinate step $h_i = h_{\text{rel}} \max(|\theta_i|, 1)$ at $h_{\text{rel}} = 10^{-5}$, symmetrised as $(F_{\text{num}} + F_{\text{num}}^T)/2$. Because the score itself is obtained by automatic differentiation, this is a single numerical derivative rather than two. That battery (data directory 433ed13) is retained in the deposit as the demonstration that the two constructions target the same information, an agreement the time-resolved figures exhibit per interval. No number reported in this paper is taken from it.

Comparing the score covariance J against the Fisher information is the sandwich, or information-matrix-equality, diagnostic (*White, 1982*): equality $\mathbf{H} = J$ holds only when the likelihood is correctly specified, and the extent to which the two differ is what the distortion matrix quantifies. The comparison has a further premise on the least-squares branch, where $\bar{G}$ presumes homoscedastic residuals; the Diagnostics section states that premise where the distortion is defined.

Uncertainty on every reported statistic comes from a nonparametric bootstrap over groups, reported as percentile intervals at the 2.5, 16, 50, 84 and 97.5% quantiles. Despite the internal column names, no probit transform is applied; the intervals are ordinary empirical percentiles, verified equal to the R quantile of type 6. The diagnostic battery used 100 bootstrap replicates and a maximum lag of 10 for the score autocorrelation, and the empirical-versus-theoretical distortion capstone used 200. The per-group estimate cloud is deposited raw. The program does run a group bootstrap of it, with a floor of 10 groups below which it falls back to a single unresampled pass, but that result is held in memory and never written: the cloud's only output file is the per-group run log, one row per group and component. Every interval quoted from the cloud in this paper is therefore computed downstream in the figure code, from the raw estimates.

Regime grid and reduction

The acquisition interval was swept over seven levels of $\tilde{\Delta}$, namely 1, 0.5, 0.2, 0.1, 0.05, 0.02 and 0.01, realised as 500, 250, 100, 50, 25, 10 and 5 raw samples averaged per interval, that is step durations of 10, 5, 2, 1, 0.5, 0.2 and 0.1 ms, giving 10, 20, 50, 100, 200, 500 and 1,000 steps in the record. The total record duration is held fixed at 100 ms across the sweep, so the interval is varied without varying the amount of data. Within each record the pre-pulse, agonist and post-pulse segments hold a fixed ratio of 20:40:40 percent of the steps.

The noise axis needs its conversion stated once, because the quantity the figures plot and the quantity the engine consumes differ by a fixed factor. The engine parameter is `Current_Noise`, a white-noise power spectral density in units of $g^2 \text{ s}$, and it reaches the likelihood only through the variance of one recorded point, $e_i = \text{Current_Noise}/\Delta_i$, with no correlation time of its own. Two design quantities are reported in the natural units of the model, and a tilde marks them: time is measured in the channel time constant $\tau = 1/k_{\text{off}}$ and current in the unitary current i . The acquisition interval is therefore $\tilde{\Delta} = \Delta k_{\text{off}} = \Delta/\tau$, and the instrumental noise is

$$\tilde{S} = \frac{S k_{\text{off}}}{i^2},$$

with S the instrumental noise power spectral density, so $\tilde{S}$ is the instrumental noise accumulated over one channel time constant measured against the single-channel signal power. It is the only

dimensionless combination of S , k_{off} and i , and it carries no free constant. At the run values $k_{\text{off}} = 100 \text{ s}^{-1}$ and $i = 1$ the reference cell used throughout is $\tilde{S} = 0.01$, and $\tilde{S} = 1$ is the level at which the instrumental noise over one channel time constant equals the single-channel signal power. The accent carries one other meaning in this paper, the boundary-pair contractions of the Theory section, Eq. 20 and Eq. 21; those are written over bold symbols and the two uses never meet in one expression.

The mean channel number N_{ch} and the noise level are the other two map axes. The results reported here are read from three data directories, each named by the git commit hash of the build that wrote it: 1c2ae6f, 87889e6 and 0ffbda7, searched in that order with the first hit winning. Every cell they contribute to a body figure is at $n_{\text{sim}} = 10^4$, which the figure code enforces by matching `n_sim_10000` in the filename, so a cell run at a smaller ensemble cannot enter a map silently. The noise axis is a fixed ladder of twelve levels, and every cell of every algorithm sits on one of them:

$$\tilde{S} \in \{0.005, 0.01, 0.02, 0.05, 0.1, 1, 10, 10^2, 10^3, 10^4, 10^5, 10^6\}, \quad (43)$$

denser below one, where the crossover the map measures sits, and by decades above it. The level $\tilde{S} = 1$ is where the instrumental noise over one channel time constant equals the single-channel signal power, so the ladder is dense on the side where the gating fluctuations still carry information and coarse on the side where they do not. Across the three directories, IR covers $N_{\text{ch}} \in \{5, 10, 20, 50, 100, 200, 500, 1000, 2000, 5000, 10000\}$ and all twelve levels; NR covers $N_{\text{ch}} \in \{10, 100, 1000, 10000\}$ and R covers those and additionally 5, 20 and 50 at $\tilde{S} = 0.01$, with the highest levels run only at the larger channel counts and reaching $\tilde{S} = 10^6$ at $N_{\text{ch}} = 10^4$; MR and VR cover the same channel numbers at $\tilde{S} \in \{0.01, 0.1, 1\}$; and the least-squares arm covers 30 cells on $N_{\text{ch}} \in \{10, 100, 1000, 10000\}$, with $\tilde{S}$ from 0.01 up to 10^3 , 10^4 , 10^5 and 10^6 respectively. The noise range is deliberately not rectangular: the highest levels were run only at the larger channel numbers, where the crossover the map measures sits. The figure code auto-detects which cells exist on disk at render time and prints the detected grid into the render log, so the cell set behind each figure is recoverable from that log.

Each grid cell is reduced through five stages: the per-group MLE cloud; the pooled joint fit θ_{pool} ; the score at both anchors, whose across-replicate covariance is J and whose per-interval Gaussian terms sum to the anchor $\tilde{G}$; the diagnostic battery at each anchor; and the empirical-versus-theoretical distortion capstone, anchored at θ_{pool} . Each cell writes five comma-separated-value families (the estimate cloud, the pooled fit, the two anchor batteries, and the capstone). The scripts number these stages one, two, four, five and six; the gap at three is the numerical-Fisher stage, which the Gaussian anchor removes and which the numbering keeps visible. On the grid-production lane the cloud stage and the capstone are dropped, because no map reads them there; the cells that need a cloud are produced by the separate script that keeps that stage.

Reproducibility and computational environment

The engine is `macro_dr`, built as C++20, and each analysis is a script in a small typed domain-specific language run in a single process (Moffatt, 2025). Provenance is self-documenting at three levels: on every invocation the binary snapshots the assembled script and its full argument list into a run ledger; every analysis output file stamps the build's short git commit hash as its first row; and the data directories are named by that hash, so results from different code versions can never overwrite one another and multi-commit provenance is documented per file.

The production batteries ran under SLURM at 32 CPUs per task and 48 GB of memory (96 GB for the parameter-recovery lane), with the linear-algebra back end pinned to a single thread so that OpenMP parallelises the per-simulation and bootstrap loops. The environment variable `MACRODR_AXIS_SERIAL=1` is load-bearing: it serialises the internal grid-combination loop, without which memory grows with the number of concurrent combinations and the jobs exhaust memory.

The one reproducibility limit worth stating twice is the simulation seed. Because the seed was the random sentinel and the resolved value was not recorded, the exact deposited ensembles

cannot be regenerated bit for bit; a rerun yields a statistically equivalent, numerically different ensemble. All resampling that operates on a fixed dataset (the bootstraps, the group partitioning) is deterministic and reproducible given the data. Deposition of the data and code, with pinned versions, is stated in the Data and code availability section.

Data and code availability

All code needed to reproduce the study is openly available. The analysis pipeline that generates the figures, together with the simulation configurations that define every run, is deposited in a public repository (`macroir-validity`) and archived at Zenodo, with the DOI issued on acceptance. The simulation and inference engine is the `macro_dr` software *Moffatt (2025)*, referenced at the pinned version that produced the deposited data; each output file records the engine commit hash that generated it. No experimental data were used: every dataset in the paper is simulated, and is regenerable from the deposited configurations.

A separate library, `macroir`, carries the boundary-conditioned likelihood, its score and its Gaussian Fisher information out of the engine and into a small self-contained C++ core, with an R package and a Python package over it. The scope of that library is narrower than the paper: it implements the boundary-conditioned member, the exact simulator and a Gauss-Newton fit, and it does not implement least squares, the open-loop member or the recursive member on instantaneous samples, so the comparisons reported here are reproducible only from `macro_dr`. The core was transcribed from the engine and checked against it on eight cells, covering four kinetic schemes, three acquisition intervals spanning a factor of a hundred, a protocol whose concentration jumps split sampling intervals, and a recording with a gap, agreeing on the total log-likelihood to between 10^{-16} and 6×10^{-15} relative and on the score to between 8×10^{-14} and 1.4×10^{-9} . The two language packages are bindings over that one core rather than independent reimplementations, and they were run against it once on a single cell, by `tools/cross_language_check.py`, which is the evidence for the statement that the three interfaces return the same numbers.

Acknowledgements

The author thanks colleagues at INQUIMAE and the University of Buenos Aires for discussion.

Author contributions

L.M. conceived the study, developed the theory, implemented the algorithms, designed and ran the simulations, analysed the data, and wrote the manuscript.

Competing interests

The author declares no competing interests.

Funding

This work was supported by the Consejo Nacional de Investigaciones Científicas y Técnicas (CONICET) and the Universidad de Buenos Aires.

Appendix 1

The interval update, member by member

The Theory section gives the cycle for the boundary-conditioned member and names the two places where the others differ. This appendix writes every member out in one template, so that each is fixed by three switches on a single set of equations.

Objects computed once per interval

All members share the quantities built from the rate matrix $\mathbf{Q}$ and the recording filter for a window of length Δ : the transition matrix $\mathbf{P}(\Delta)$; the boundary-conditioned single-channel moments $\bar{\gamma}_{i \rightarrow j}$ and $\bar{v}_{i \rightarrow j}$; the start-conditioned mean conductance $\bar{\gamma}_0$ of Eq. 10; the start-conditioned interval variance $\mathbf{w}$; and, for the members that ignore the acquisition averaging, the instantaneous per-state conductance vector γ . Two matrices collect the boundary weights,

$$G_{i \rightarrow j} = \bar{\gamma}_{i \rightarrow j} P_{i \rightarrow j}(\Delta), \quad H_{i \rightarrow j} = \bar{\gamma}_{i \rightarrow j}^2 P_{i \rightarrow j}(\Delta). \quad (\text{A1})$$

The tilde operator

Over a bold symbol, the tilde marks a contraction taken on the K^2 boundary pairs ($i \rightarrow j$) rather than on the K states. (Over a scalar it means something else, a quantity in the natural units of the model, as in $\tilde{\Delta}$ and $\tilde{S}$; the two uses are disjoint and never meet in one expression.) It is not a new operator: its two cases are the two contractions the update of Eq. 11 already has, $\gamma^\top \Sigma \gamma$ and $\Sigma \gamma$, evaluated there with the boundary objects substituted. Both are written out in the Theory section, the bilinear one as the scalar of Eq. 20 and the vector one as the gain $\mathbf{g}_{0 \rightarrow \Delta} = \widetilde{\gamma^\top \Sigma}$ of Eq. 21, and both collapse to $K \times K$ arithmetic once the boundary covariance Eq. 15 is substituted, so the $K^2 \times K^2$ object is never built. The matrices $\mathbf{G}$ and $\mathbf{H}$ of Eq. A1 are what tells the two apart, one power of $\bar{\gamma}$ against two.

One cycle, three switches

Every member on the lattice runs the same four steps on each window: propagate the occupancy Gaussian (Eqs. 16 and 17), predict the observation (Eqs. 18 and 19), condition on the recorded average (Eqs. 22 and 23), and either carry the posterior to the next window or discard it. Three switches fix which member is being run.

Switch 1: the window.

It has three settings, an instant, the state at the window's start, and the states at both of its ends, and it selects the conductance object, and with it the second term of the predictive variance and the gain. At the instantaneous setting the observation is treated as a sample taken at the centre of the window: the occupancy is propagated there first, by $\mathbf{P}(\Delta/2)$, the conductance is γ , and, with Σ read in that mid-window frame,

$$\widetilde{\gamma^\top \Sigma \gamma} \rightarrow \gamma^\top \Sigma \gamma, \quad \mathbf{g} \rightarrow \mathbf{P}(\Delta/2)^\top \Sigma \gamma. \quad (\text{A2})$$

At the start-conditioned setting the same two expressions hold with $\bar{\gamma}_0$ in place of γ and the full $\mathbf{P}(\Delta)$ in the gain, since the conductance is conditioned on the state at the window's start and the whole window remains to be crossed; the contraction is still an ordinary quadratic form on the K states. The trailing propagator in the gain is not decoration. It carries the update from the frame in which the observation was formed to the frame the next window starts in, which is the end of the current window, so it is the remaining half at the instantaneous setting and the whole window at the start-conditioned one. Omitting it, or using the full window where half is called for, leaves the rank-one down-date in the wrong frame, and

the symptom is specific and was observed: a spike in the distortion-induced bias of the recursive instantaneous member at long intervals and large channel counts. At the boundary setting the conductance is conditioned on both endpoints and the two expressions become the tildes, Eq. 20 and Eq. 21. The step from the start state to the boundary therefore does two things at once: it replaces Σ_0 by $\Sigma_0 - \text{diag}(\mu_0)$ and adds the boundary term ($\mu_0^T \mathbf{H} \mathbf{1}$ in the variance, $\mathbf{G}^T \mu_0$ in the gain). The subtracted diagonal is not a second approximation: the within-channel part it removes is already inside the boundary term, and dropping the subtraction would count it twice.

Switch 2: recursion.

A recursive member carries $(\mu^{\text{post}}, \Sigma^{\text{post}})$ into the next window. A non-recursive one discards it and propagates the next window's predictive from the initial condition, so its log-likelihood factorizes over windows. The switch changes no formula above; it changes only what is fed to the next one.

Switch 3: the interval variance, total or residual.

It selects the third term of Eq. 19, which is present only when the acquisition averaging is modelled at all, that is at either of the two interval settings of switch 1:

$$(\mathbf{w})_i = \begin{cases} \sum_j P_{i \rightarrow j}(\Delta) [\bar{v}_{i \rightarrow j} + (\bar{y}_{i \rightarrow j} - \bar{y}_{i \rightarrow \cdot})^2] & \text{total,} \\ \sum_j P_{i \rightarrow j}(\Delta) \bar{v}_{i \rightarrow j} & \text{residual.} \end{cases} \quad (\text{A3})$$

The two differ by the spread of the interval mean across end states, which is Eq. 26. The switch is free only where the window is conditioned on the start state alone. With both endpoints conditioned on, the residual form is forced whatever the flag says, for the same reason the diagonal is subtracted in switch 1: the between-endpoint spread is already carried by the boundary term, so the total form would count it twice.

The members

Member	window	recursive	interval variance	what it is
NR	an instant	no	—	instantaneous samples, windows independent
R	an instant	yes	—	instantaneous samples, posterior carried forward
INR	the start	no	total	NR with the window average restored
MR	the start	yes	total	R with the window average restored
VR	the start	yes	residual	MR with switch 3 flipped, gain unchanged
IR	both ends	yes	residual (forced)	VR with the boundary terms restored

Appendix 1—table 1. The six members of the lattice as settings of the three switches. Reading down the last column, each row differs from an earlier one in exactly one switch, which is what makes the ladder measurable one step at a time. IR is MacroIR; INR is the MacroINR of *Moffatt and Pierdominici-Sottile (2025)*. The data keys under which each is dispatched are in Table 2.

One consequence of the two switches acting together is worth stating, because it is easy to read the table as saying otherwise. Evaluated at the *same* prior, MR and IR assign the

interval average the same total variance. Expanding the total form of Eq. A3 and using $\sum_j P_{i \rightarrow j}(\Delta) = 1$ and $\sum_j P_{i \rightarrow j}(\Delta) \tilde{y}_{i \rightarrow j} = \tilde{y}_{i \rightarrow \cdot}$

$$\mu_0^T \mathbf{w} \Big|_{\text{total}} = \mu_0^T \mathbf{w} \Big|_{\text{residual}} + \mu_0^T \mathbf{H} \mathbf{1} - \tilde{\mathbf{y}}_0^T \text{diag}(\mu_0) \tilde{\mathbf{y}}_0, \quad (\text{A4})$$

and the two extra pieces are exactly what switch 1 moves from the third term of Eq. 19 into the second when the window goes from the start state to both ends. The decompositions differ, the sum does not. What does change between the two members is the gain, and at a common prior that is the whole of the MR-to-IR difference. It does not follow that the two report the same variance on a recording. The predicted variance divides the gain, so a member's variance sets the state the next window opens with: the moment the gains differ the priors differ, and every quantity downstream of the prior, the predicted variance included, differs with them. The identity above is between the two maps from a prior to a predicted variance, and it is tested only where the two members still hold the same prior.

Two of the six pairs isolate one algebraic change at the point where it is made, and they are the ones the Results use. Stepping MR to VR flips switch 3, and stepping VR to IR restores the boundary terms in both the variance and the gain while leaving the variance form where it already was. Neither step is confined to what it flips, and the reason is the same one as above: the predictive variance divides the gain, so flipping switch 3 moves the update as well, and the two steps are separable in the algebra without being separable in what a recording does with them. The two together are the R-to-IR gap.

Least squares, which is off the lattice

Classical nonlinear least squares has no setting of the three switches. It propagates the mean occupancy alone, $\mu_\Delta = \mathbf{P}(\Delta)^T \mu_0$, with no covariance and no update, predicts $\tilde{y}_{0 \rightarrow \Delta}^{\text{pred}}$ from Eq. 18, and assigns one constant variance to every residual, fitted as the residual sum of squares over the number of samples. Neither term of Eq. 19 that depends on the gating survives, and there is no gain, so no occupancy object is ever conditioned on the data. It is the anchor of the comparison rather than a member of the family.

References

- Anderson CR**, Stevens CF. Voltage Clamp Analysis of Acetylcholine Produced End-Plate Current Fluctuations at Frog Neuromuscular Junction. *The Journal of Physiology*. 1973; 235(3):655–691. doi: [10.1113/jphysiol.1973.sp010410](https://doi.org/10.1113/jphysiol.1973.sp010410).
- Celentano JJ**, Hawkes AG. Use of the Covariance Matrix in Directly Fitting Kinetic Parameters: Application to GABAA Receptors. *Biophysical Journal*. 2004; 87(1):276–294. doi: [10.1529/biophysj.103.036632](https://doi.org/10.1529/biophysj.103.036632).
- Clerx M**, Beattie KA, Gavaghan DJ, Mirams GR. Four Ways to Fit an Ion Channel Model. *Biophysical Journal*. 2019; 117(12):2420–2437. doi: [10.1016/j.bpj.2019.08.001](https://doi.org/10.1016/j.bpj.2019.08.001).
- Colquhoun D**, Hawkes AG. On the Stochastic Properties of Single Ion Channels. *Proceedings of the Royal Society of London Series B*. 1981; 211(1183):205–235. doi: [10.1098/rspb.1981.0003](https://doi.org/10.1098/rspb.1981.0003).
- Conti F**, Neumcke B, Nonner W, Stämpfli R. Conductance Fluctuations from the Inactivation Process of Sodium Channels in Myelinated Nerve Fibres. *The Journal of Physiology*. 1980; 308:217–239. doi: [10.1113/jphysiol.1980.sp013279](https://doi.org/10.1113/jphysiol.1980.sp013279).
- Del Core L**, Mirams GR. Parameter Inference for Stochastic Reaction Models of Ion-Channel Gating from Whole-Cell Voltage-Clamp Data. *Philosophical Transactions of the Royal Society A*. 2025; 383(2292):20240224. doi: [10.1098/rsta.2024.0224](https://doi.org/10.1098/rsta.2024.0224).
- Godambe VP**. An Optimum Property of Regular Maximum Likelihood Estimation. *The Annals of Mathematical Statistics*. 1960; 31(4):1208–1211. doi: [10.1214/aoms/1177705693](https://doi.org/10.1214/aoms/1177705693).
- Hodgkin AL**, Huxley AF. A Quantitative Description of Membrane Current and its Application to Conduction and Excitation in Nerve. *The Journal of Physiology*. 1952; 117(4):500–544. doi: [10.1113/jphysiol.1952.sp004764](https://doi.org/10.1113/jphysiol.1952.sp004764).

- Huber PJ.** The Behavior of Maximum Likelihood Estimates under Nonstandard Conditions. In: *Proceedings of the Fifth Berkeley Symposium on Mathematical Statistics and Probability, Vol. 1* Berkeley: University of California Press; 1967.p. 221–233. <https://projecteuclid.org/euclid.bsmmsp/1200512988>.
- Katz B, Miledi R.** Membrane Noise Produced by Acetylcholine. *Nature*. 1970; 226:962–963. doi: 10.1038/226962a0.
- Katz B, Miledi R.** The Statistical Nature of the Acetylcholine Potential and its Molecular Components. *The Journal of Physiology*. 1972; 224(3):665–699. doi: 10.1113/jphysiol.1972.sp009918.
- Lei CL, Ghosh S, Whittaker DG, Aboelkassem Y, Beattie KA, Cantwell CD, Delhaas T, Houston C, Novaes GM, Panfilov AV, Pathmanathan P, Riabiz M, dos Santos RW, Walmsley J, Carr H, Gavaghan DJ, Mirams GR.** Considering Discrepancy when Calibrating a Mechanistic Electrophysiology Model. *Philosophical Transactions of the Royal Society A*. 2020; 378(2173):20190349. doi: 10.1098/rsta.2019.0349.
- Milescu LS, Akk G, Sachs F.** Maximum Likelihood Estimation of Ion Channel Kinetics from Macroscopic Currents. *Biophysical Journal*. 2005; 88(4):2494–2515. doi: 10.1529/biophysj.104.053256.
- Moffatt L.** Estimation of Ion Channel Kinetics from Fluctuations of Macroscopic Currents. *Biophysical Journal*. 2007; 93(1):74–91. doi: 10.1529/biophysj.106.101212.
- Moffatt L.** MacrolR: Bayesian Kinetic Modeling of P2X2 Activation. Zenodo; 2025. <https://doi.org/10.5281/zenodo.17162475>, doi: 10.5281/zenodo.17162475, software, version v1.0.0.
- Moffatt L, Hume RI.** Responses of Rat P2X2 Receptors to Ultrashort Pulses of ATP Provide Insights into ATP Binding and Channel Gating. *Journal of General Physiology*. 2007; 130(2):183–201. doi: 10.1085/jgp.200709779.
- Moffatt L, Pierdominici-Sottile G.** Bayesian Inference of Functional Asymmetry in the P2X2 Receptor. *Communications Biology*. 2025; 8. doi: 10.1038/s42003-025-09056-x.
- Münch JL, Paul F, Schmauder R, Benndorf K.** Bayesian Inference of Kinetic Schemes for Ion Channels by Kalman Filtering. *eLife*. 2022; 11:e62714. doi: 10.7554/eLife.62714.
- Neher E, Stevens CF.** Conductance Fluctuations and Ionic Pores in Membranes. *Annual Review of Biophysics and Bioengineering*. 1977; 6:345–381. doi: 10.1146/annurev.bb.06.060177.002021.
- Owen MJ, Mirams GR.** IonBench: A Benchmark of Optimisation Strategies for Mathematical Models of Ion Channel Currents. *PLoS Computational Biology*. 2025; 21(8):e1013319. doi: 10.1371/journal.pcbi.1013319.
- Sigworth FJ.** Covariance of Nonstationary Sodium Current Fluctuations at the Node of Ranvier. *Biophysical Journal*. 1981; 34(1):111–133. doi: 10.1016/S0006-3495(81)84840-0.
- Stepanyuk A, Borisyuk A, Belan P.** Maximum Likelihood Estimation of Biophysical Parameters of Synaptic Receptors from Macroscopic Currents. *Frontiers in Cellular Neuroscience*. 2014; 8:303. doi: 10.3389/fn-cel.2014.00303.
- Stepanyuk AR, Borisyuk AL, Belan PV.** Efficient Maximum Likelihood Estimation of Kinetic Rate Constants from Macroscopic Currents. *PLoS ONE*. 2011; 6(12):e29731. doi: 10.1371/journal.pone.0029731.
- Van Loan CF.** Computing Integrals Involving the Matrix Exponential. *IEEE Transactions on Automatic Control*. 1978; 23(3):395–404. doi: 10.1109/TAC.1978.1101743.
- White H.** Maximum Likelihood Estimation of Misspecified Models. *Econometrica*. 1982; 50(1):1–25. doi: 10.2307/1912526.
- Zadrozny PA.** Gaussian Likelihood of Continuous-Time ARMAX Models When Data Are Stocks and Flows at Different Frequencies. *Econometric Theory*. 1988; 4(1):108–124. doi: 10.1017/S0266466600011890.




Figure 1—figure supplement 1. The same filter cycle over the first ten milliseconds of the recording, so that the concentration jump and the onset of gating are inside the window.




Figure 3—figure supplement 1. Per-step information against score variance, the opening rate and the closing rate. Seven members across the columns, the least-squares arm being the interval-averaged ILSE, and two parameters down the page as bands of two rows: k_{on} (A, B) and k_{off} (C, D). The upper row of each band carries the per-interval Fisher information and the score variance as magnitudes, which lie on each other where the member is calibrated at that interval; the lower row is $\log_{10}(J_t/F_t)$ with its 95% bootstrap interval, zero when they do. Per-interval calibration is not monotone along the cost ladder: on k_{off} the median reads +0.184 for ILSE, -0.038 for NR, +0.007 for INR, -0.145 for R, -0.229 for MR, -0.135 for VR and +0.033 for IR, so the cheapest member that models the interval average is better calibrated interval by interval than any of the three recursive members that cost more. Curves stop at the information mask, which is also where the upper axis ends. In the two recursive columns the k_{on} band reaches that floor within six intervals of the agonist being removed while the k_{off} band does not. The least-squares ratio row is the same curve in every band of both halves, with a median of 0.183889 in each: with a single constant variance the parameter cancels out of the ratio, which leaves the true residual variance over the assumed one.




Figure 3—figure supplement 2. Per-step information against score variance, the unitary current and the channel number. The other half of the preceding supplement, on the same seven members and read the same way: the unitary current (A, B) and N_{ch} (C, D), the upper row of each band the per-interval Fisher information against the score variance and the lower row $\log_{10}(J_i/F_i)$ with its 95% bootstrap interval. In the two recursive columns the N_{ch} band reaches the information floor within six intervals of the agonist being removed while the unitary-current band does not, which is the contrast the Discussion draws on how a recording is spent. The channel-number band carries no least-squares panel, which with the gating fluctuations unmodelled cannot separate that parameter from the unitary current.




Figure 3—figure supplement 3. The checks this figure has no room for. The seven members of the preceding supplements at the same cell, carrying (A) the mean standardized residual, the one of the three properties of r_t , not drawn above; (B) the score bias for the three parameters not coloured above, k_{on} , the unitary current and the instrumental noise level; and (C) their score autocorrelation. The first residual moment separates nobody, which is what makes the other two properties the informative ones. In (B) the ordering repeats the one in figure supplement 2: counting informative intervals whose interval excludes zero against a chance level of 0.05, IR reads 0.067, 0.050 and 0.050, INR 0.100, 0.100 and 0.000, and the three recursive members between 0.409 and 0.975. The noise level is the control direction: it is the only parameter measured before the agonist arrives, and its score is white even in the members whose kinetic directions are not, so the memory that carries the accumulated information error lives in the gating directions and not in the instrumental one. Least squares carries no noise curve because its noise level is a plug-in from the same record rather than a free direction, so its score there vanishes identically, which is the same fact as its mean squared residual being pinned at one.




Figure 4—figure supplement 1. The distortion decomposed into a per-sample part and a correlation part, against channel count and resolved by instrumental noise. The per-sample parts of the two recursive members coincide and go to one as channels are added, while the correlation parts separate.




Figure 4—figure supplement 2. The sandwich prediction against the empirical multivariate distribution of the estimates, as a Mahalanobis quantile-quantile plot with a χ^2 reference. The reported Fisher covariance under-covers the nominal 95% region and the sandwich covers it.




Figure 4—figure supplement 3. The size of the information distortion against its shape. Two invariants of the same distortion matrix over the same plane, for six members. (A) the geometric mean of its eigenvalues, the typical factor by which the reported information is wrong over all directions at once, signed about a null of one. (B) the spread of those eigenvalues about that mean, how much the factor varies from direction to direction, one-sided about a null of one. The two are the orthogonal components of the affine-invariant distance, so (B) at one means a single effective-count rescaling is exact and (A) is the factor it should use, and (B) above one means no such rescaling exists at any factor.